



逢愛君 / Lecturer

馮才昇 李峻宇 劉厚辰 / T.A.s


SIMULATING SDN WITH MININET

Outline

- SDN Overview
- Mininet Overview
- Ryu Overview
- Lab Announcement
- Mininet VM & Ryu tutorial
- OpenFlow Protocol and Open vSwitch
- Reference

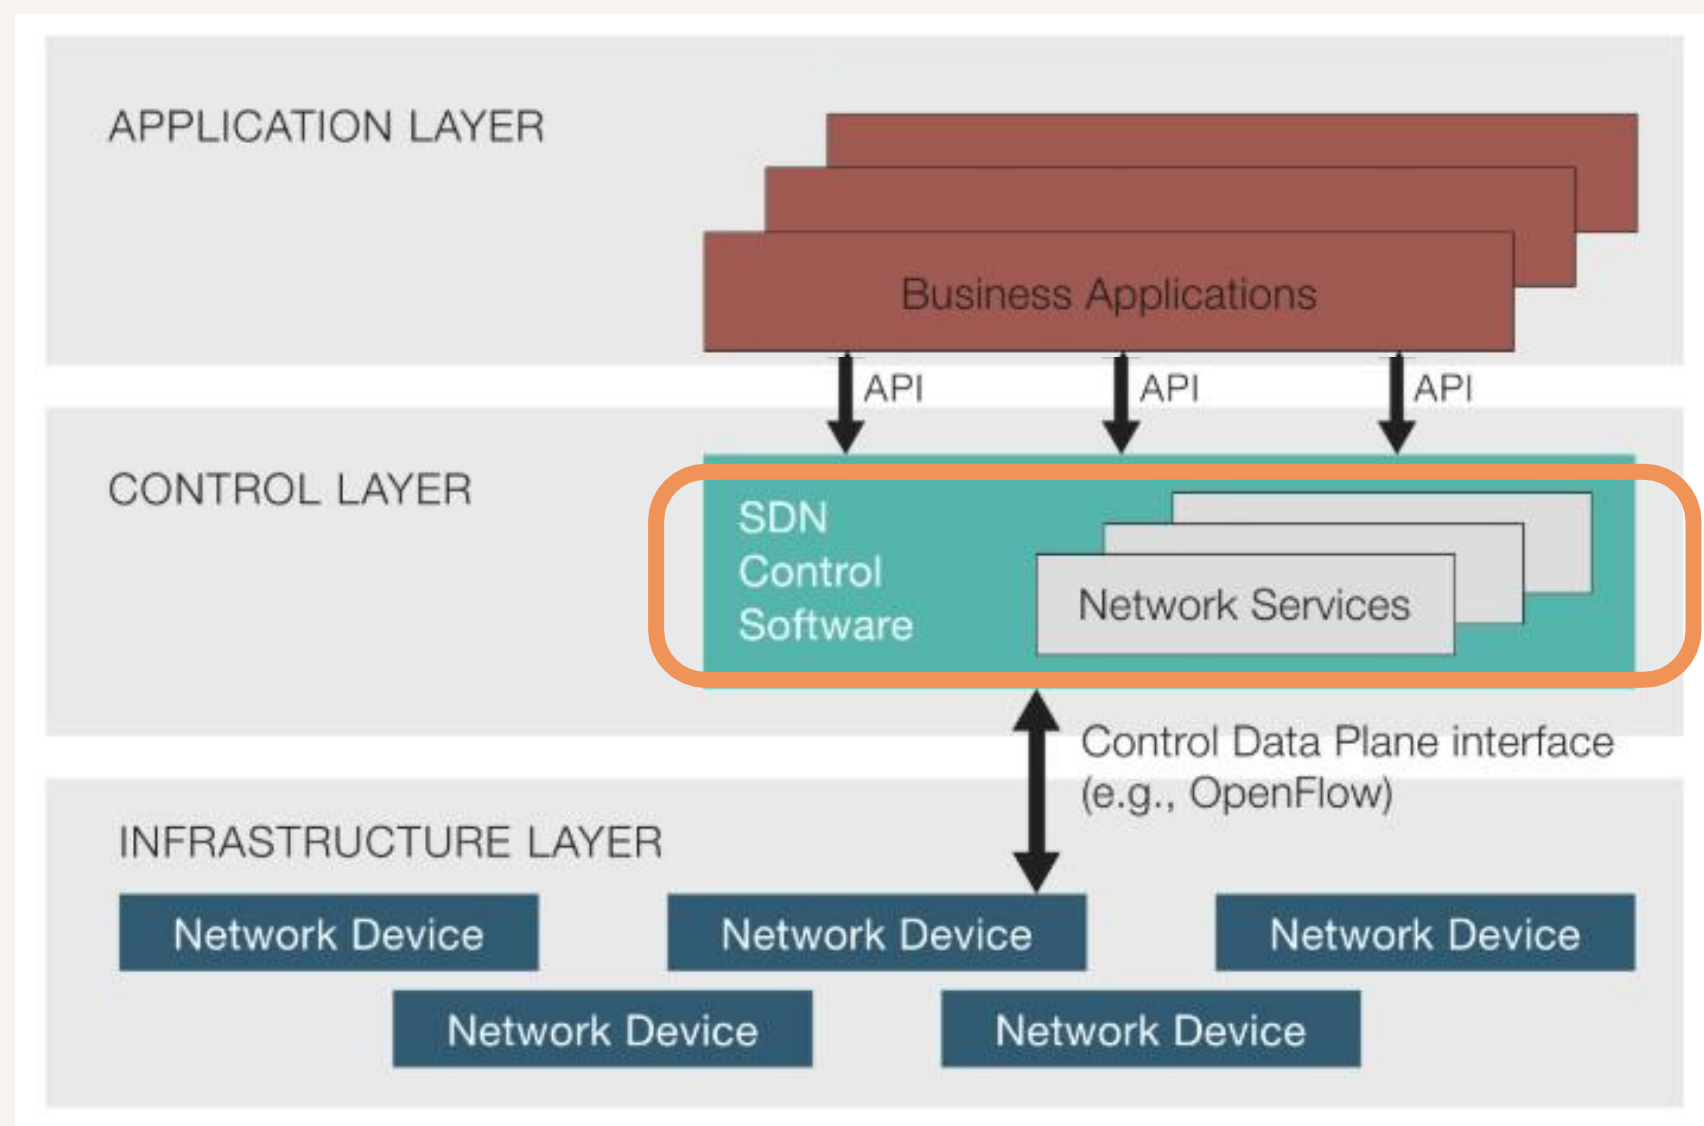
SDN Overview

- Decoupling of control and data planes
- Directly Programmable network
- Logically centralized network



SDN Overview

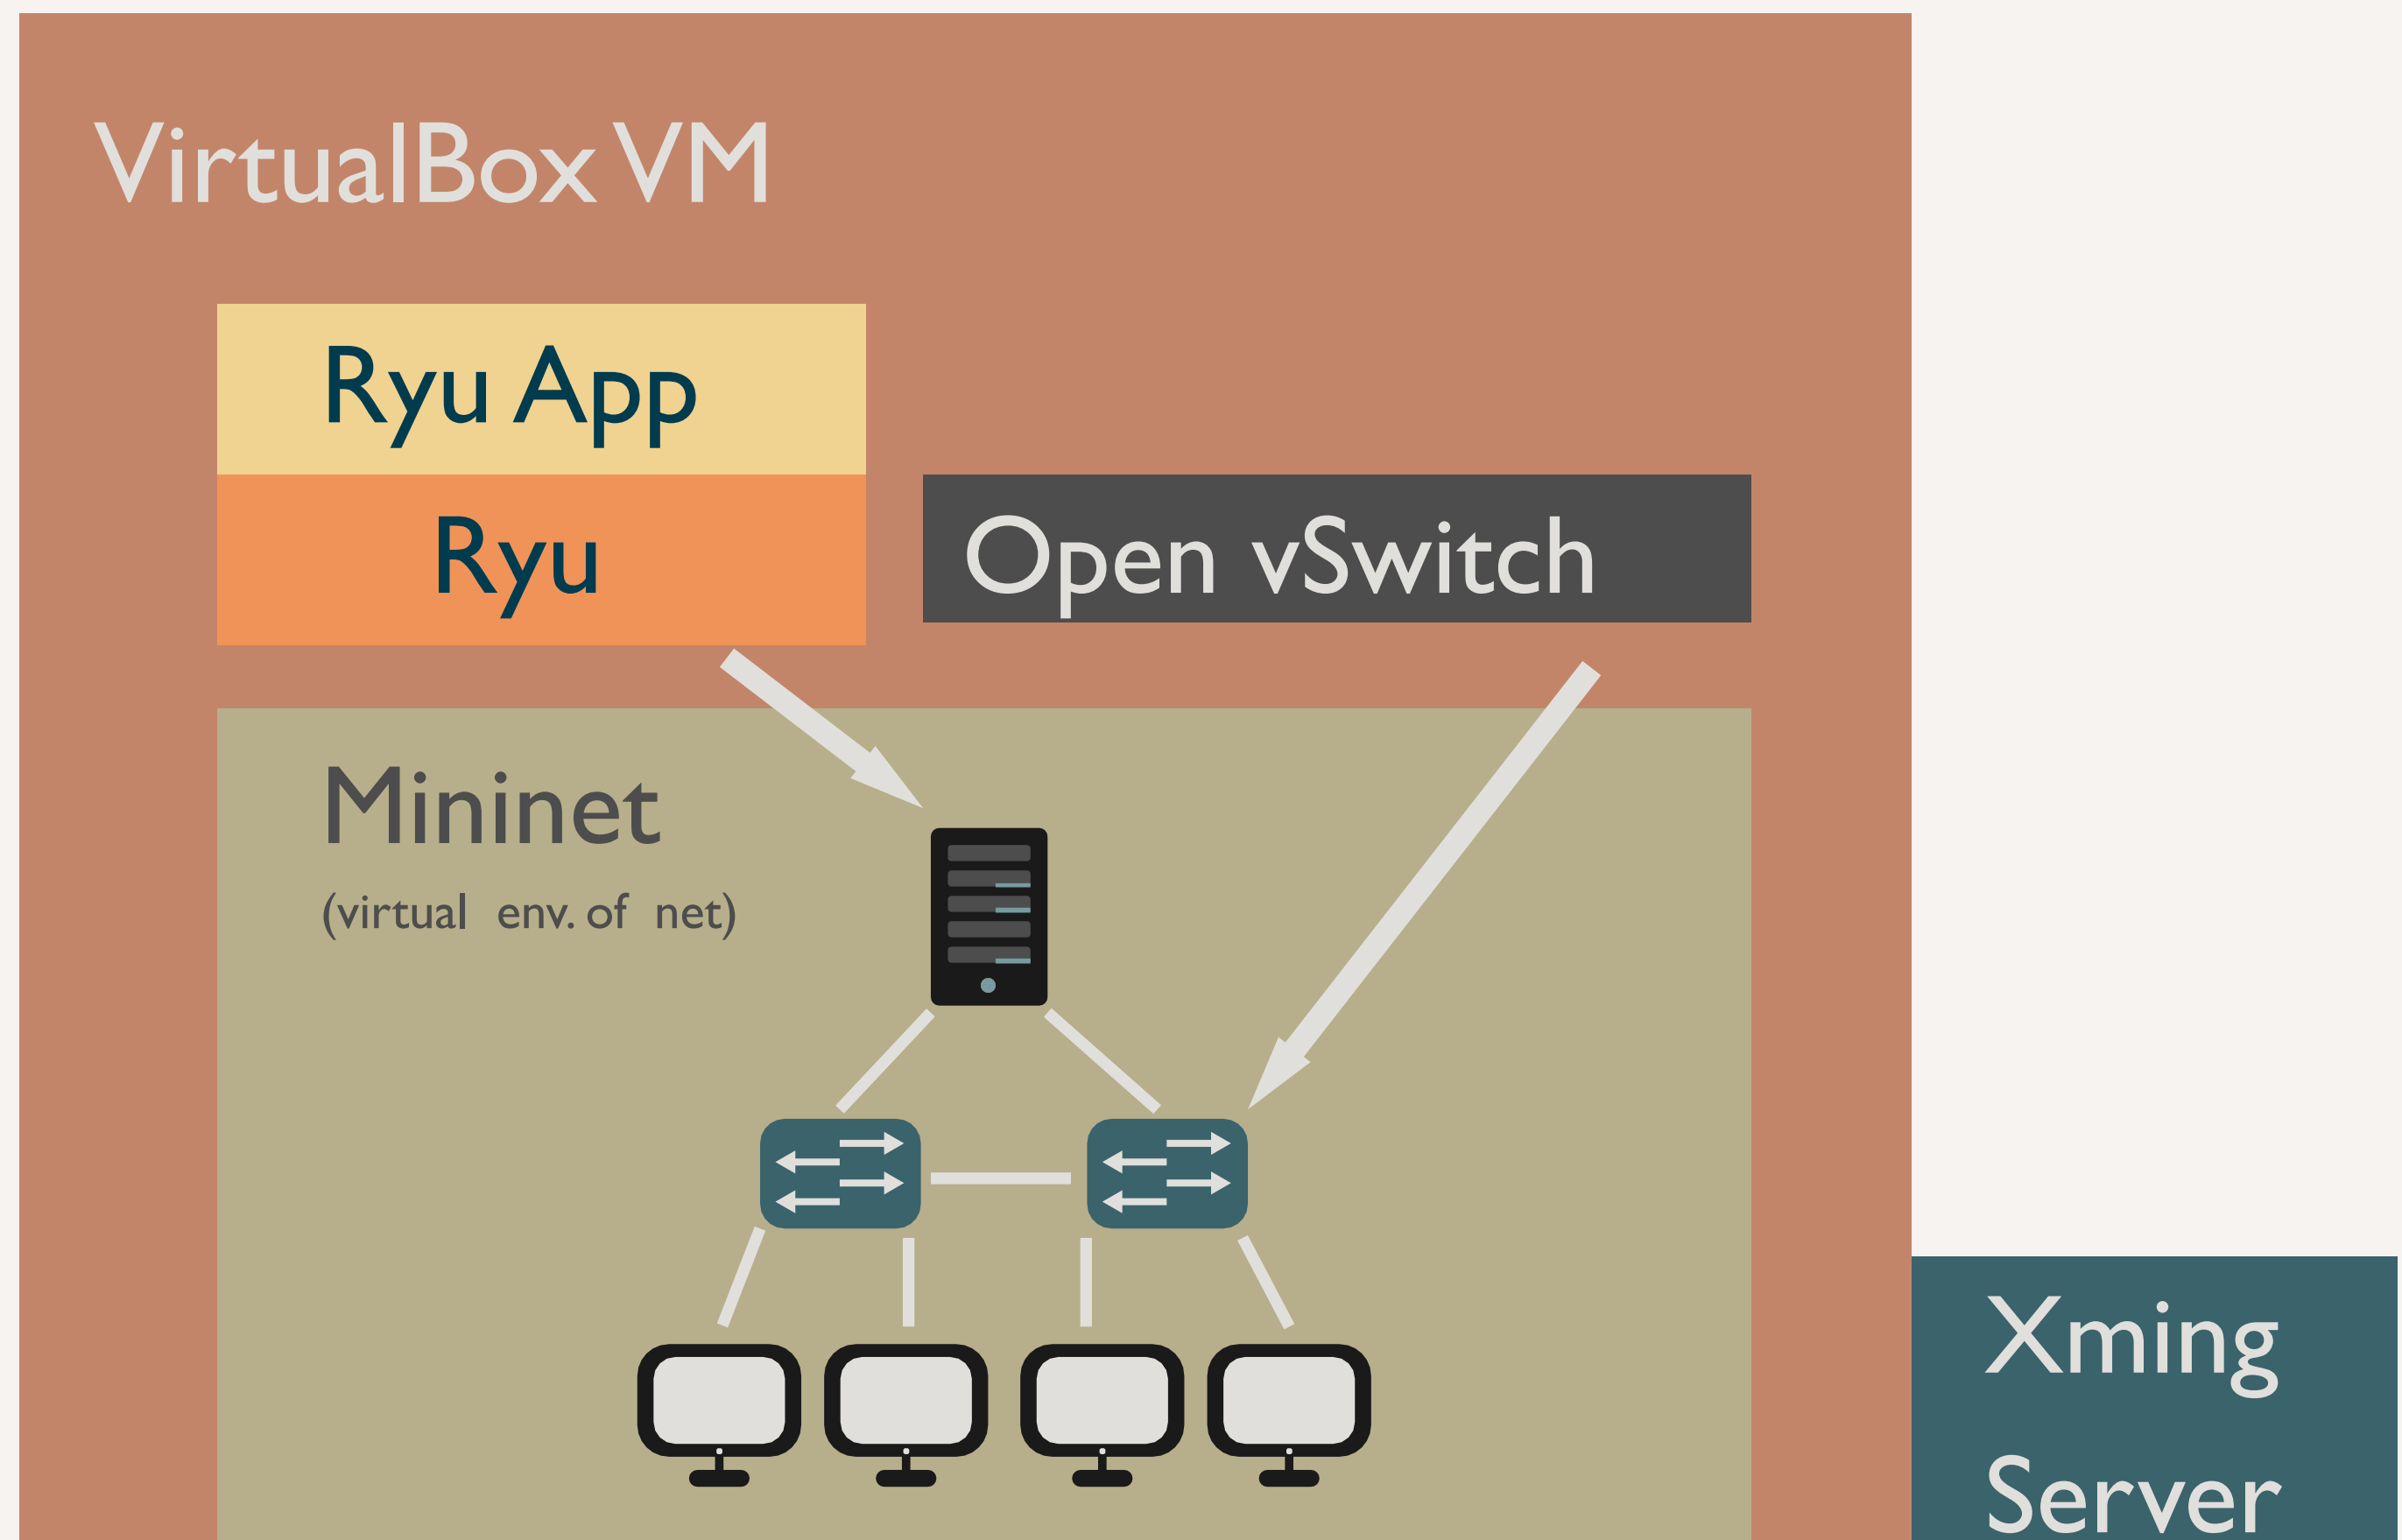
- Under SDN architecture, Ryu framework is in the control layer.
- It helps developers to communicate between the controller and underlying Open vSwitch.



Ryu Framework

**Mininet and
Open vSwitch**


Mininet and Ryu Overview





Mininet

- A **network emulator** that create virtual network topology running real kernel, switch or application.
- Includes a **CLI** (**C**ommand **L**ine **I**nterface), for debugging or running network-wide tests.
- Supports arbitrary **custom topologies**, and includes a basic set of parametrized topologies(Linear, Tree).
- Provides a straightforward and extensible **Python API** for network creation and experimentation.

Network Emulator



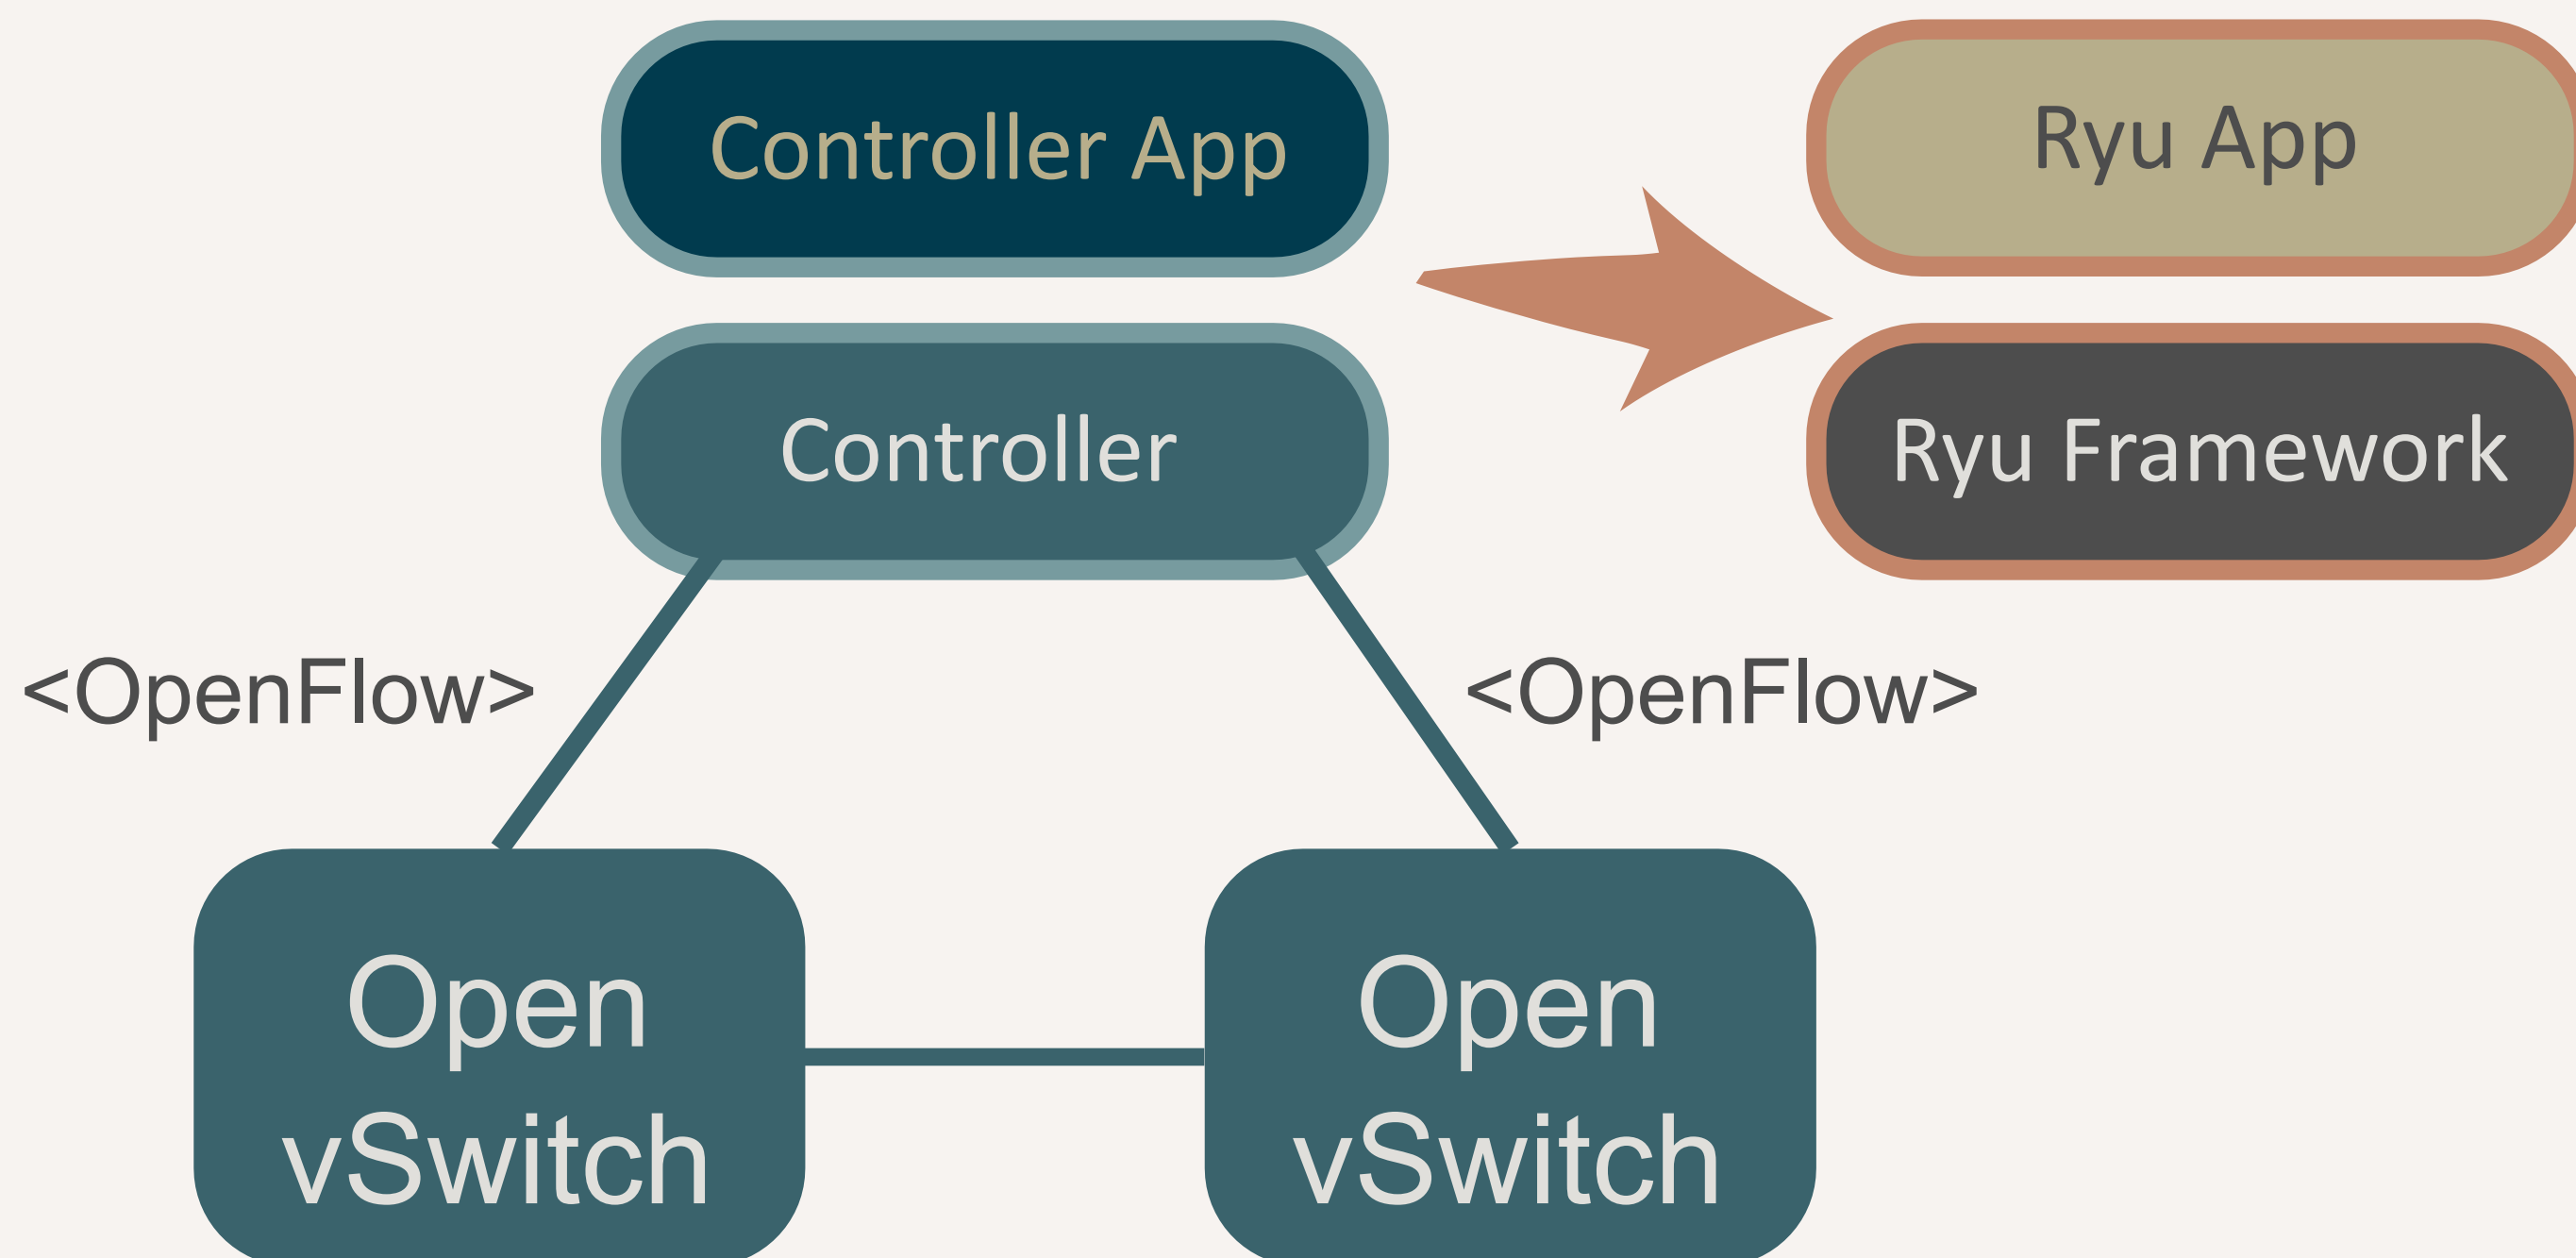
Emulated
Network



Hardware
Network

Ryu Overview

- Ryu can help developers to communicate between the controller and Open vSwitch by using OpenFlow protocol.




Lab Announcement


Specification

- Create your network topology by Mininet
 - A tree topology
 - A data center-like topology
- Write your SDN application by Ryu
 - Simple network slicing
 - Network monitor
 - Congestion detection
 - Congestion reaction, drop or re-routing
- Report

Mininet – Tree Topology



Mininet – Data Center-Like Topology



Mininet –Data Center-Like Topology


- Your mininet program should be able to receive input which specifies the number of core and aggregation switches.
- There are loops in the topology, which will cause **broadcast storm**. How to solve it?

Ryu application – Simple Network Slicing

- Network slicing is a way to share infrastructure, it overlays multiple virtual network on top of a shared network.
- Controller have the overall view of the physical network, and can install rules to “slice” the network.
- Hosts or switches in different will not know each other and therefore can not access to each other.


Ryu application – Simple Network Slicing

Topology from
controller's
view




Ryu application – Simple Network Slicing

Network 1



Ryu application – Simple Network Slicing

Network 2



Ryu application – Simple Network Slicing

- To simplify your work, you can slice the network by IP address only(e.g. class 1 contains 10.0.0.<odd number>, class 2 contains 10.0.0.<even number>)
- In this task, you can use destination MAC address to transfer packets and flood packet to unknown hosts(learning switch).

Ryu application – Network Monitor

You need to periodically send request to switch and receive port and flow statistical information from switches.

- Port statistical information
 - send/receive packet count
 - send/receive byte count
 - drop count
 - error count
- Flow statistical information
 - in-port
 - match fields
 - action
 - match packet/byte count

Ryu application – Network Monitor

- You need to print out the statistical information received by your application.
- Below is a example format for printing out, you can use arbitrary format to show the information
- How to define a flow?

Flow Statistical Information											
datapath	in-port	match	src_ip	src_port	dst_ip	dst_port	protocol	action	packets	bytes	
0000000000000001	2	10.0.0.2	33868	10.0.0.1	5001	TCP	dropped	4	296		
0000000000000001	2	10.0.0.2	33872	10.0.0.1	5001	TCP	dropped	3	222		
0000000000000001	3	10.0.0.3	58212	10.0.0.1	5001	TCP	port 1	106	1055572		
0000000000000001	1	10.0.0.1	5001	10.0.0.3	58212	TCP	port 3	81	5346		
Port Statistical Information											
datapath	port	rx-pkts	rx-bytes	tx-pkts	tx-bytes	dropped	error				
0000000000000001	1	97	6454	158	1060678	0	0				
0000000000000001	2	22	1616	50	4990	0	0				
0000000000000001	3	150	1060100	116	9068	0	0				
0000000000000001	fffffffe	0	0	0	0	24	0				

Ryu application – Congestion Detection


- Base on your network monitor, now you can detection a congestion on a port of a switch
- You can calculate the delta value by the information received this time and previous time
- Once a congestion occurred(e.g. 1MB/s), your application should detect it and print out a alert message

Ryu application – Congestion Reaction

- In addition to detecting a congestion when occurred, now you should react to this congestion, you can either
 - Drop the big flow
 - Re-route the big flow
- In this task, you must identify a flow by 5-tuple field
 - source IP address(e.g. 10.0.0.1)
 - source port(e.g. 5001)
 - destination IP address(e.g. 10.0.0.3)
 - destination port(e.g. 33872)
 - Protocol(e.g. TCP, UDP)


Ryu application – Congestion Example

- This is the path from h1 to h2 originally. We can send packets from h1 to h2.



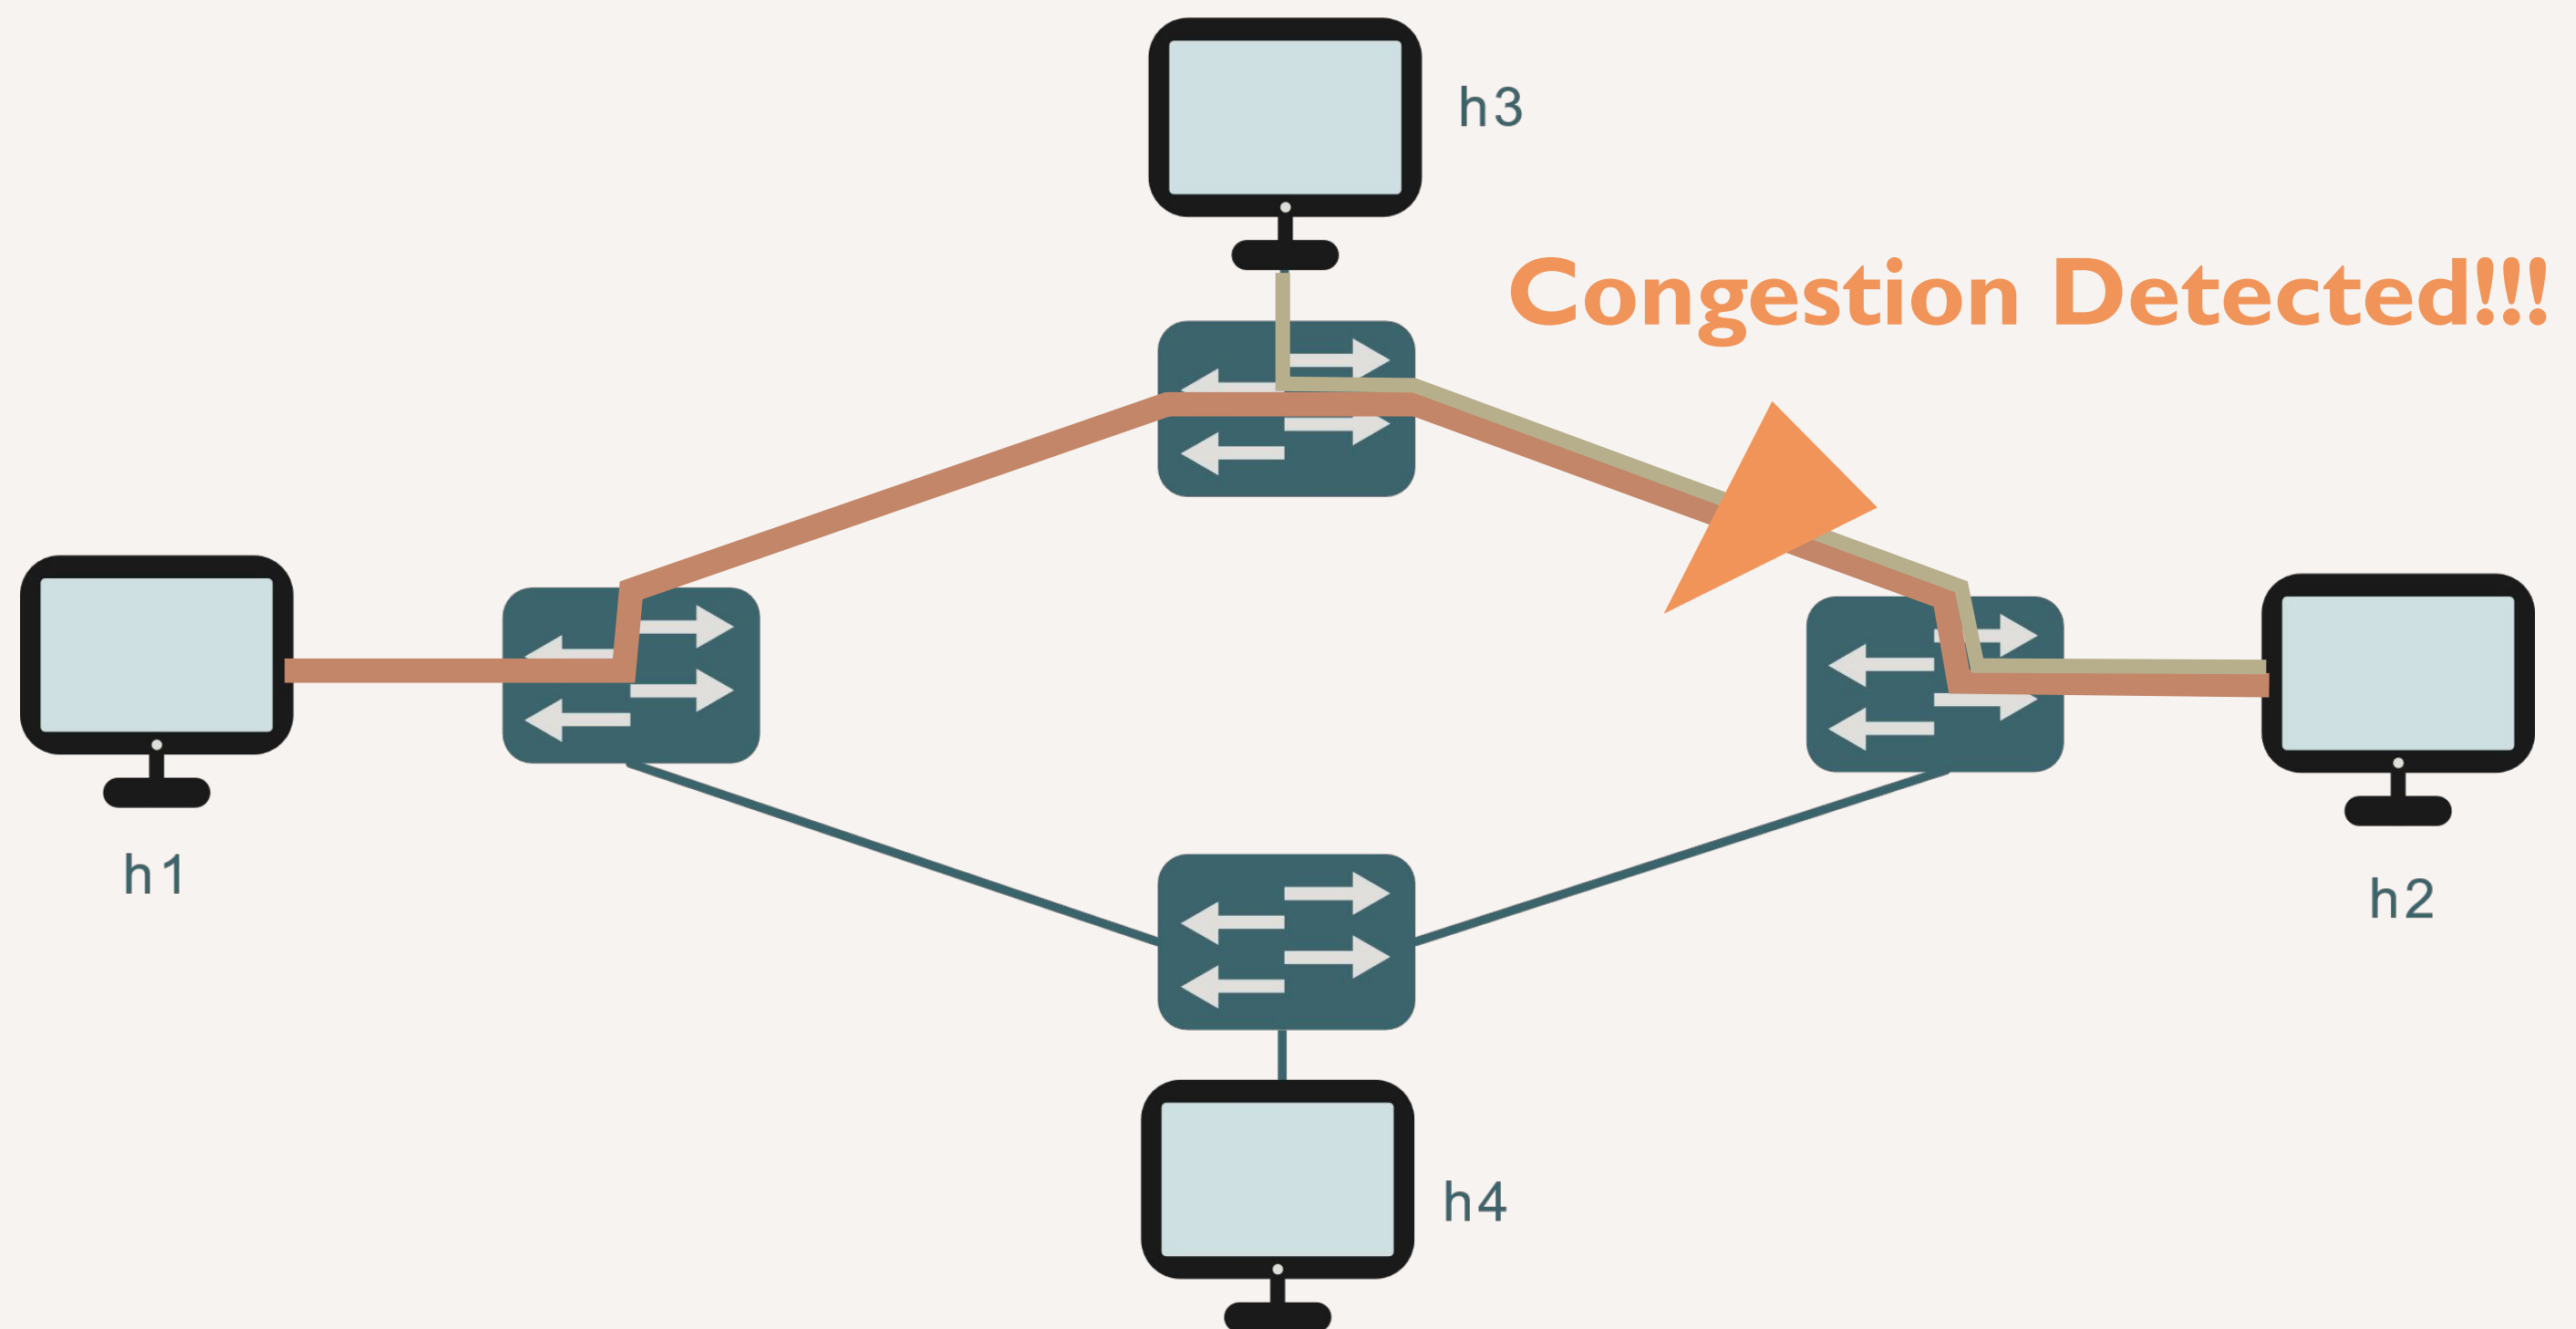
Ryu application – Congestion Example

- However, there may be others packets on the path from h1 to h2.




Ryu application – Congestion Example

- While congestion occurs, your application should be able to detect it. (20%)




Ryu application – Congestion Example

- Your application can drop the big flow.



Ryu application – Congestion Example

- Your application can re-route the big flow.



Grading Policy

- **Mininet** **(25%)**
 - A tree topology (15%)
 - A data center like topology (10%)
- **Ryu application** **(50%)**
 - Simple network slicing (8% + 2%)
 - Network monitor (8% + 2%)
 - Congestion detection (5% + 5%)
 - Congestion reaction - drop (5% + 5%)
 - Congestion reaction - re-route (10%)
- **Report** **(25%)**

Report

- **Environment (5%)**
- **Explain the pros and cons that there are loops in a network topology. (10%)**
- **Explain the broadcast storm and how you handle it in this lab. Is there any better solution to handle broadcast storms under SDN? If yes, explain how; if no, explain why. (10%)**

Deadline

- **5 / 27 Demo**
- **6 / 3 23:59:59**
 - Submit report and source code to NTU cool
 - Report has to be in pdf format

Mininet

Mininet mn Linux Commands

- This command needs **superuser privilege(sudo)**.
- Basic CLI

Command Line	Description
\$ sudo mn	Start a minimal topology
\$ exit	Exit mininet
\$ sudo mn -c	Exit and free the resource

Mininet CLI Commands

- Display Mininet CLI commands

```
$ mininet> help
```

- Display nodes

```
$ mininet> nodes
```

- Display links

```
$ mininet> net
```

- Dump information about all nodes

```
$ mininet> dump
```

Mininet mn Linux Commands

- `-h, --help` # Show the help message.
- `--controller` # Define what controller to use.
- `--topo` # Tell mininet what topology to use.
- `--custom` # Read custom topology and node
parameter from .py file.
- `-x, --xterm` # Spawn an xterm window for all
nodes inside the network.
- `--mac` # The mininet will assign static MAC
addr. for each host.

Mininet CLI Commands

- Xterm display

```
$ mininet> xterm h1
```

- Ping between all hosts

```
$ mininet> pingall
```

- <node> <command>

-Ex1: \$ mininet> h1 ifconfig

-Ex2: \$ mininet> h1 ping h2

```
mininet> h1 ifconfig
h1-eth0  Link encap:Ethernet  HWaddr f6:db:51:50:c1:f1
         inet addr:10.0.0.1  Bcast:10.255.255.255  Mask:255.0.0.0
         UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
         RX packets:0 errors:0 dropped:0 overruns:0 frame:0
         TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
         collisions:0 txqueuelen:1000
         RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

lo       Link encap:Local Loopback
         inet addr:127.0.0.1  Mask:255.0.0.0
         UP LOOPBACK RUNNING  MTU:65536  Metric:1
         RX packets:0 errors:0 dropped:0 overruns:0 frame:0
         TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
         collisions:0 txqueuelen:0
         RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

mininet>
```

views on the mininet from host1

Xterm on Mininet CLI

```
$ sudo mn --topo single,3
```

```
$ mininet> xterm h1
```

```
[2019-01-22 21:21.16] ~
[p1974.DESKTOP-HVG9A9V] > ssh mininet@192.168.56.255
ssh: connect to host 192.168.56.255 port 22: Cannot assign requested address

[2019-01-22 21:22.01] ~
[p1974.DESKTOP-HVG9A9V] > ssh mininet@192.168.56.101
Warning: Permanently added '192.168.56.101' (RSA) to the list of known hosts.
Welcome to Ubuntu 14.04.4 LTS (GNU/Linux 4.2.0-27-generic i686)

 * Documentation:  https://help.ubuntu.com/
Last login: Tue Jan 22 19:18:58 2019
mininet@mininet-vm:~$ ls
install-mininet-vm.sh  loxigen  mininet  oflops  oftest  openflow  pox
mininet@mininet-vm:~$ sudo mn --topo single,3
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1) (h3, s1)
*** Configuring hosts
h1 h2 h3
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> xterm h1
mininet> █
```

```

"Node: h1"@mininet-vm
root@mininet-vm:~# ifconfig
h1-eth0  Link encap:Ethernet  HWaddr 0e:5c:95:4c:c8:4b
         inet addr:10.0.0.1  Bcast:10.255.255.255  Mask:255.0.0.0
         UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
         RX packets:0 errors:0 dropped:0 overruns:0 frame:0
         TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
         collisions:0 txqueuelen:1000
         RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

lo       Link encap:Local Loopback
         inet addr:127.0.0.1  Mask:255.0.0.0
         UP LOOPBACK RUNNING  MTU:65536  Metric:1
         RX packets:693 errors:0 dropped:0 overruns:0 frame:0
         TX packets:693 errors:0 dropped:0 overruns:0 carrier:0
         collisions:0 txqueuelen:0
         RX bytes:1164676 (1.1 MB)  TX bytes:1164676 (1.1 MB)

root@mininet-vm:~# █
```


views on the mininet from host1
in Xterm

Mininet Create Custom Topology

- Use python script to create custom mininet topology.
- Some useful function
 - `addHost(<host name>)` // add host
 - `addSwitch(<host name>)` // add switch
 - `addLink(<node1>,<node2>)` // add link

Mininet Create Custom Topology

- Example
Topology



Mininet Create Custom Topology

```
from mininet.topo import Topo
```


```
class MyTopo (Topo):
    # "Simple topology example."
```

```
    def __init__(self):
        # "Create custom topo."
        # Initialize topology
        Topo.__init__(self)
```

```
    # Add hosts and switches
    leftHost = self.addHost('h1')
    rightHost = self.addHost('h2')
    leftSwitch = self.addSwitch('s1')
    rightSwitch = self.addSwitch('s2')
```

```
    # Add links
    self.addLink(leftHost, leftSwitch)
    self.addLink(leftSwitch, rightSwitch)
    self.addLink(rightSwitch, rightHost)
```

```
topos = {'mytopo': (lambda: MyTopo())}
```



→ Setup Part

→ Create Topology Part

→ Critical Part, which names your topology

Mininet Create Custom Topology

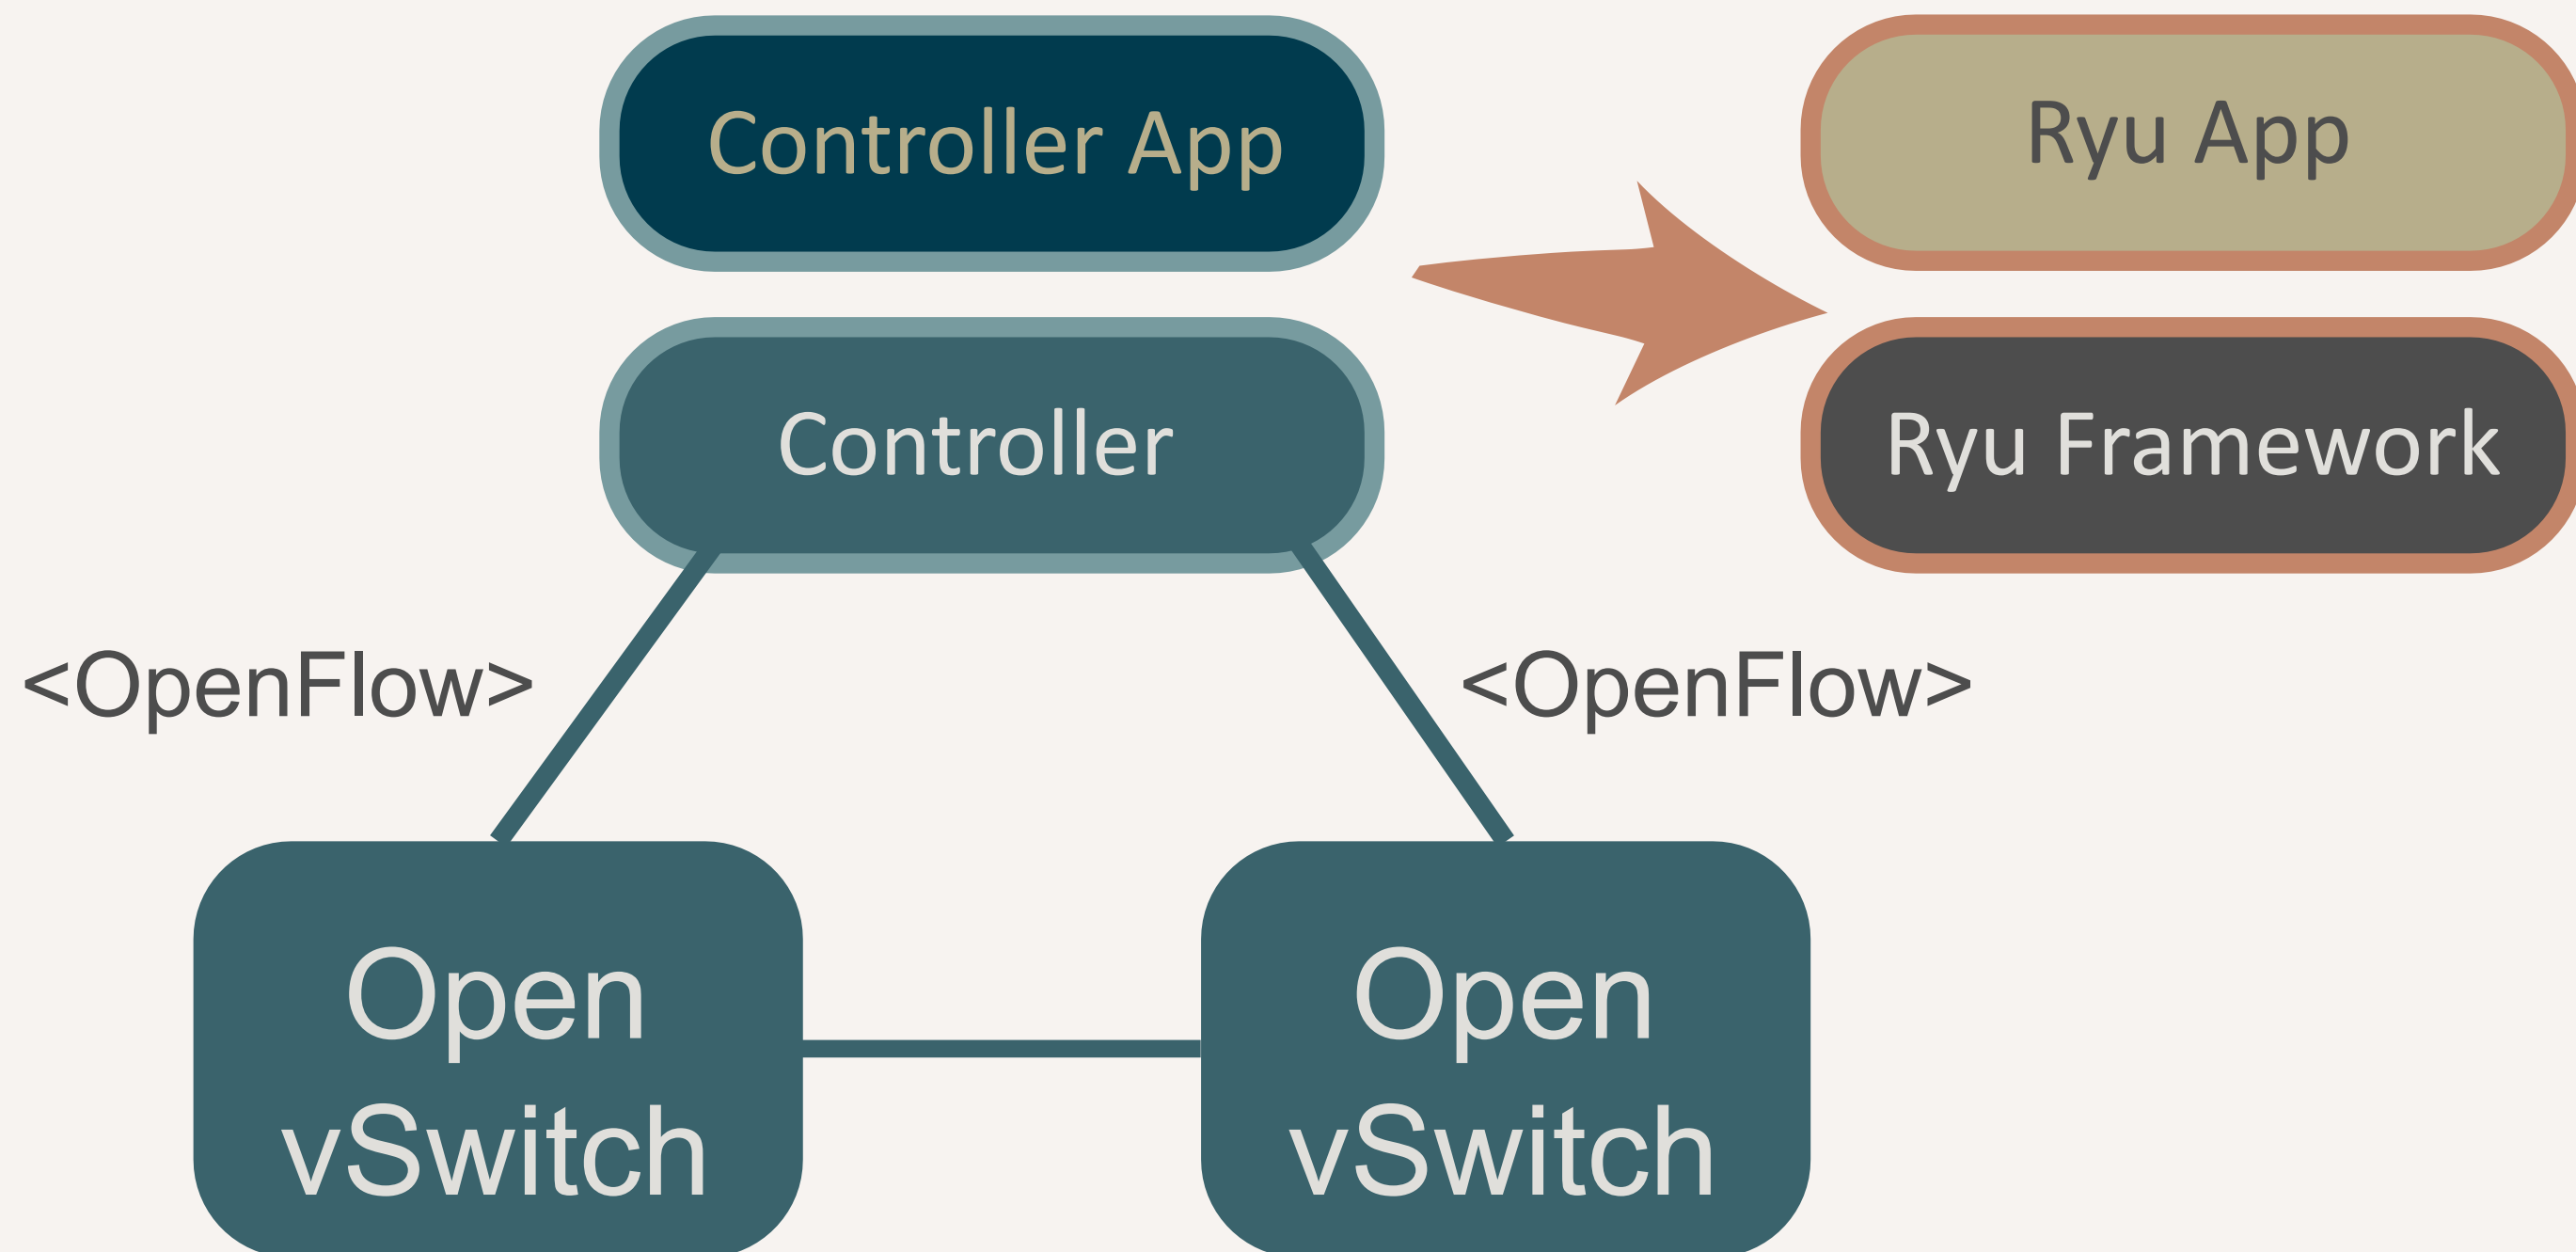
- Save the topology file in the file system
- Use the custom flag and topo flag to tell mininet what topology you want to use.

```
$ sudo mn --custom <path>/mytopology.py  
--topo=mytopo  
--controller=remote,ip=<host addr.>,port=<port>
```

↓
optional

Ryu

Recall the Conception of Ryu



Installation of Ryu

```
$ sudo apt-get update
$ sudo apt-get install python-pip python-dev -y
$ sudo apt-get install python-eventlet python-routes \
                        python-webob python-paramiko -y
$ sudo pip install ryu
$ sudo pip install --upgrade six
```

Ryu L2 Learning Switch Application

- You can start with the source code of the application ([link](#))

- Features:

- **Learning table**: MAC address to port
- **Event handler**

- Run the Mininet Topology

```
$ sudo mn --topo single,3 --mac --switch=ovsk,\nprotocols=OpenFlow13 --controller remote
```

- Run the Ryu Application

(Login VM with another xterm windows or Ctrl+Alt+F2 to tty2)

```
$ ryu-manager --verbose ryu.app.simple_switch_13
```

Ryu Class Definition

- Must inherit the “app_manager.RyuApp”

```
class SimpleSwitch13(app_manager.RyuApp):  
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]  
  
    def __init__(self, *args, **kwargs):  
        super(SimpleSwitch13, self).__init__(*args, **kwargs)  
        self.mac_to_port = {}
```

OpenFlow Messages- PacketIn

- A way for the switch to send a captured packet to the controller.
- Occurs on **two circumstances**:
 - **Base on some rule** from switch.
 - From a **miss** in the match tables, so the switch send a PacketIn by **default**.

OpenFlow Messages- PacketOut

- The **response of a PacketIn** from SDN controller
- The controller has the ability to inject packets into the data plane of a particular switch by PacketOut message, which can either
 - **carry a raw packet** to inject into the switch, or
 - indicate a **local buffer** on the switch containing a raw packet to release.

OpenFlow Messages


- FeatureRes

The FeatureRes is the switch's reply to the controller
enumerating its abilities.

- FlowMod

- One of the main messages, it allows the controller to
modify the state of an OpenFlow switch.

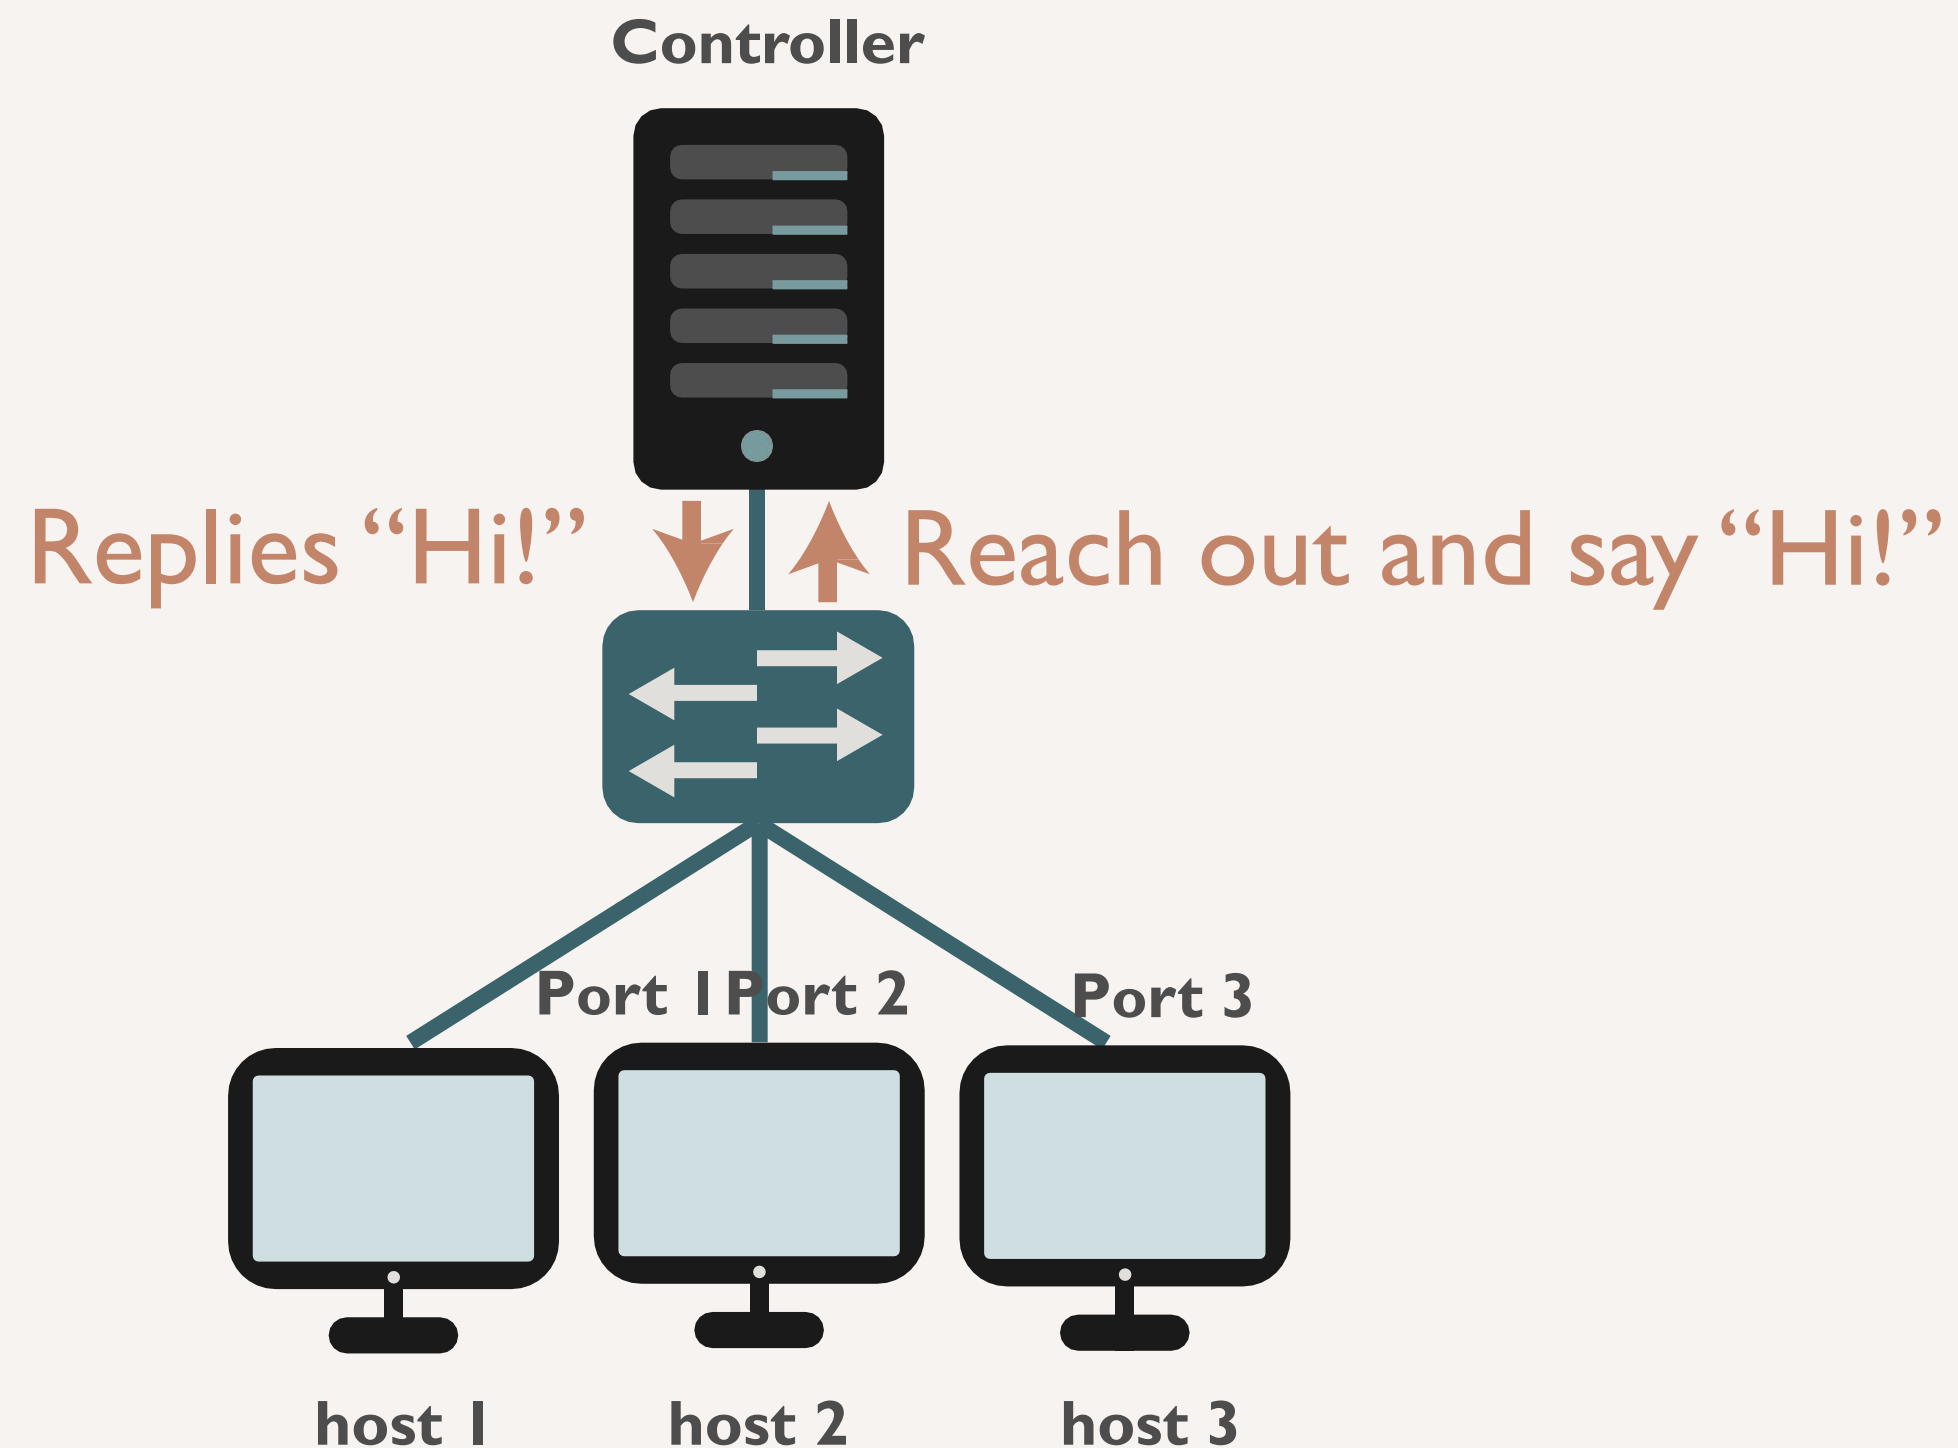
An Example of the Process



An Example of the Process

– Step 1: Handshaking


- Controller received hello message from the switch and then made the response. This has been handled by Ryu framework.



An Example of the Process

– Step 2:Featuring

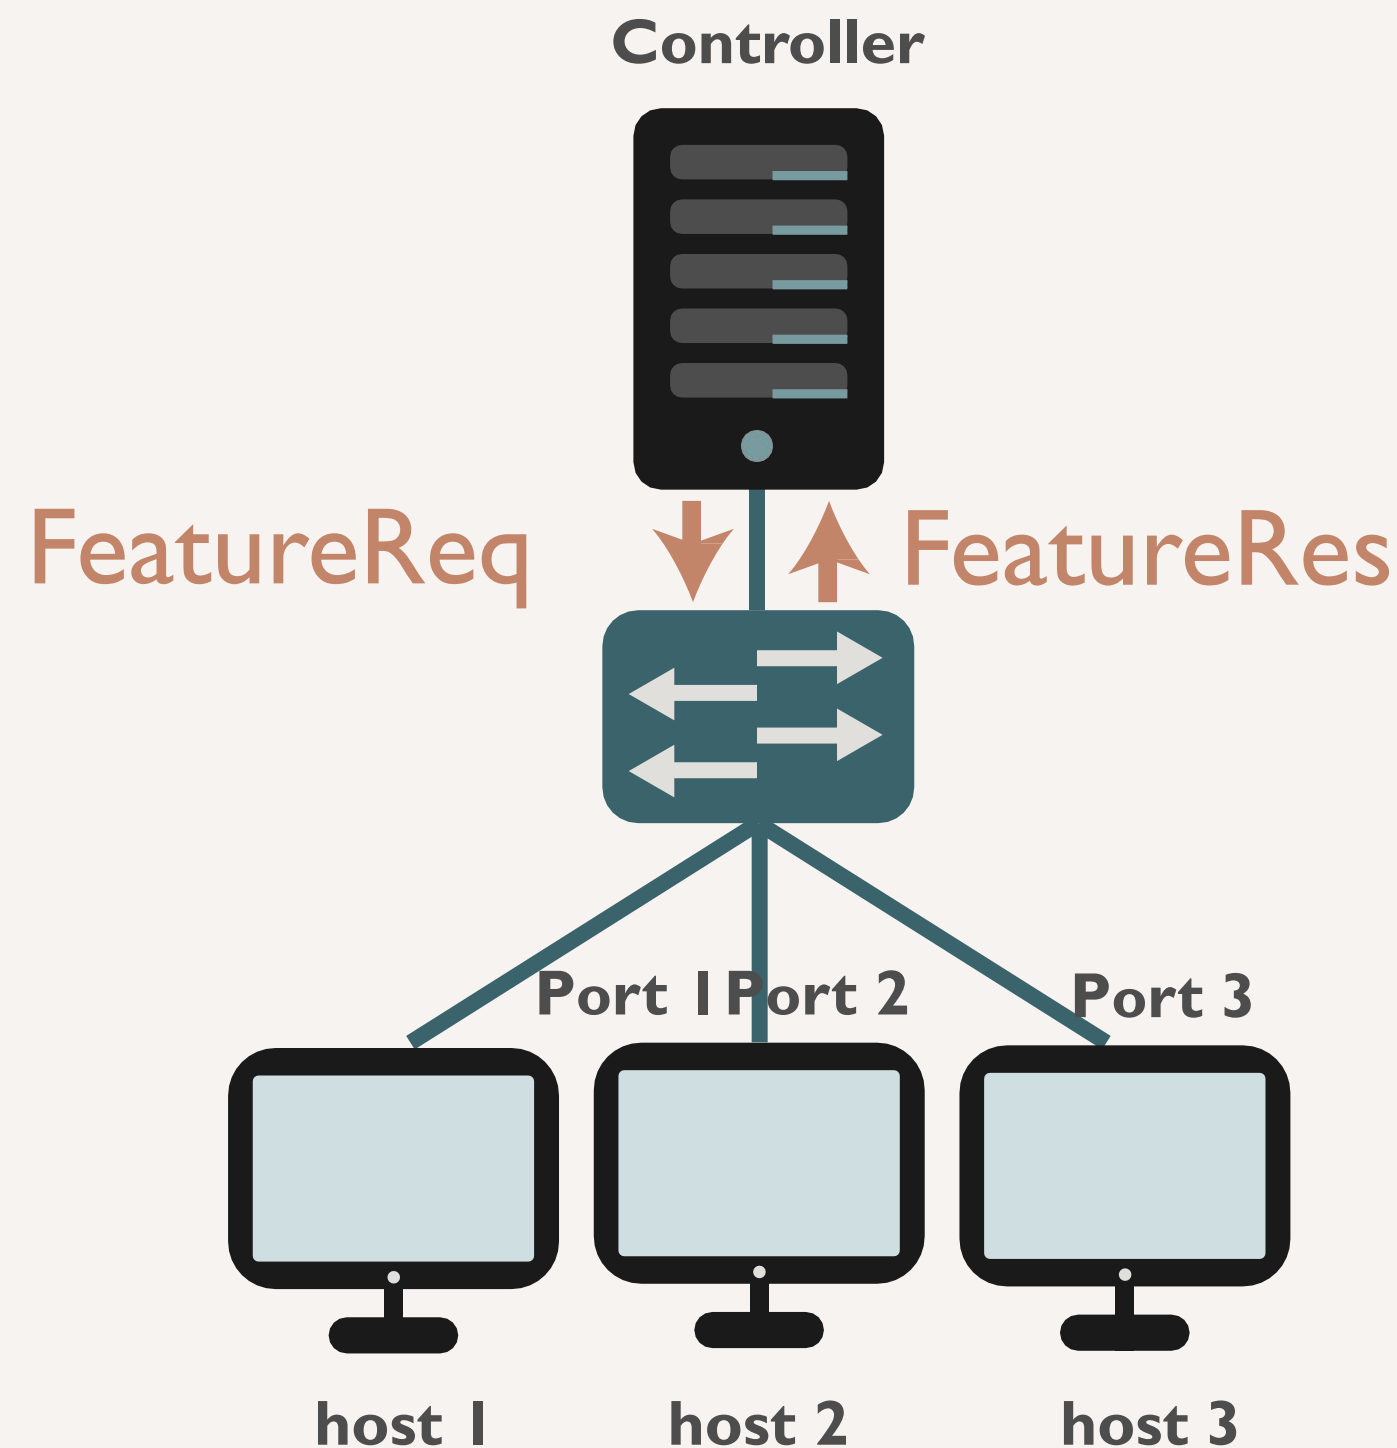
- Controller sent feature request to the switch to get its capabilities.



An Example of the Process

– Step 2:Featuring

- Switch sent **feature response** to the controller and then controller add the **lowest priority rule** to the switch.
(**table miss entry**: go to the controller)



An Example of the Process

– Step 2:Featuring

- Switch sent **feature response** to the controller and then controller add the **lowest priority rule** to the switch.

```
@set_ev_cls(ofp_event.EventOFPSwitchFeatures,
CONFIG_DISPATCHER)
def switch_features_handler(self, ev):
    datapath = ev.msg.datapath
    ofproto = datapath.ofproto
    parser = datapath.ofproto_parser
    # install table-miss flow entry
    #
    # We specify NO BUFFER to max_len of the output action due to
    # OVS bug. At this moment, if we specify a lesser number, e.g.,
    # 128, OVS will send Packet-In with invalid buffer id and
    # truncated packet data. In that case, we cannot output packets
    # correctly. The bug has been fixed in OVS v2.1.0.


    match = parser.OFPMatch()
    actions = [parser.OFPACTIONOutput(ofproto.OFPP_CONTROLLER,
                                      ofproto.OFPCML_NO_BUFFER)]
    self.add_flow(datapath, 0, match, actions)
```

#....

An Example of the Process

– Step 3: ARP table miss


- When a host (host 1) ping another host (host 3), it will trigger a table miss on Network card if the packet hasn't been seen before.



An Example of the Process

– Step 3: ARP table miss


- Then host 1 will broadcast an **ARP request** to get the MAC address of host 3.



An Example of the Process

– Step 3: ARP table miss

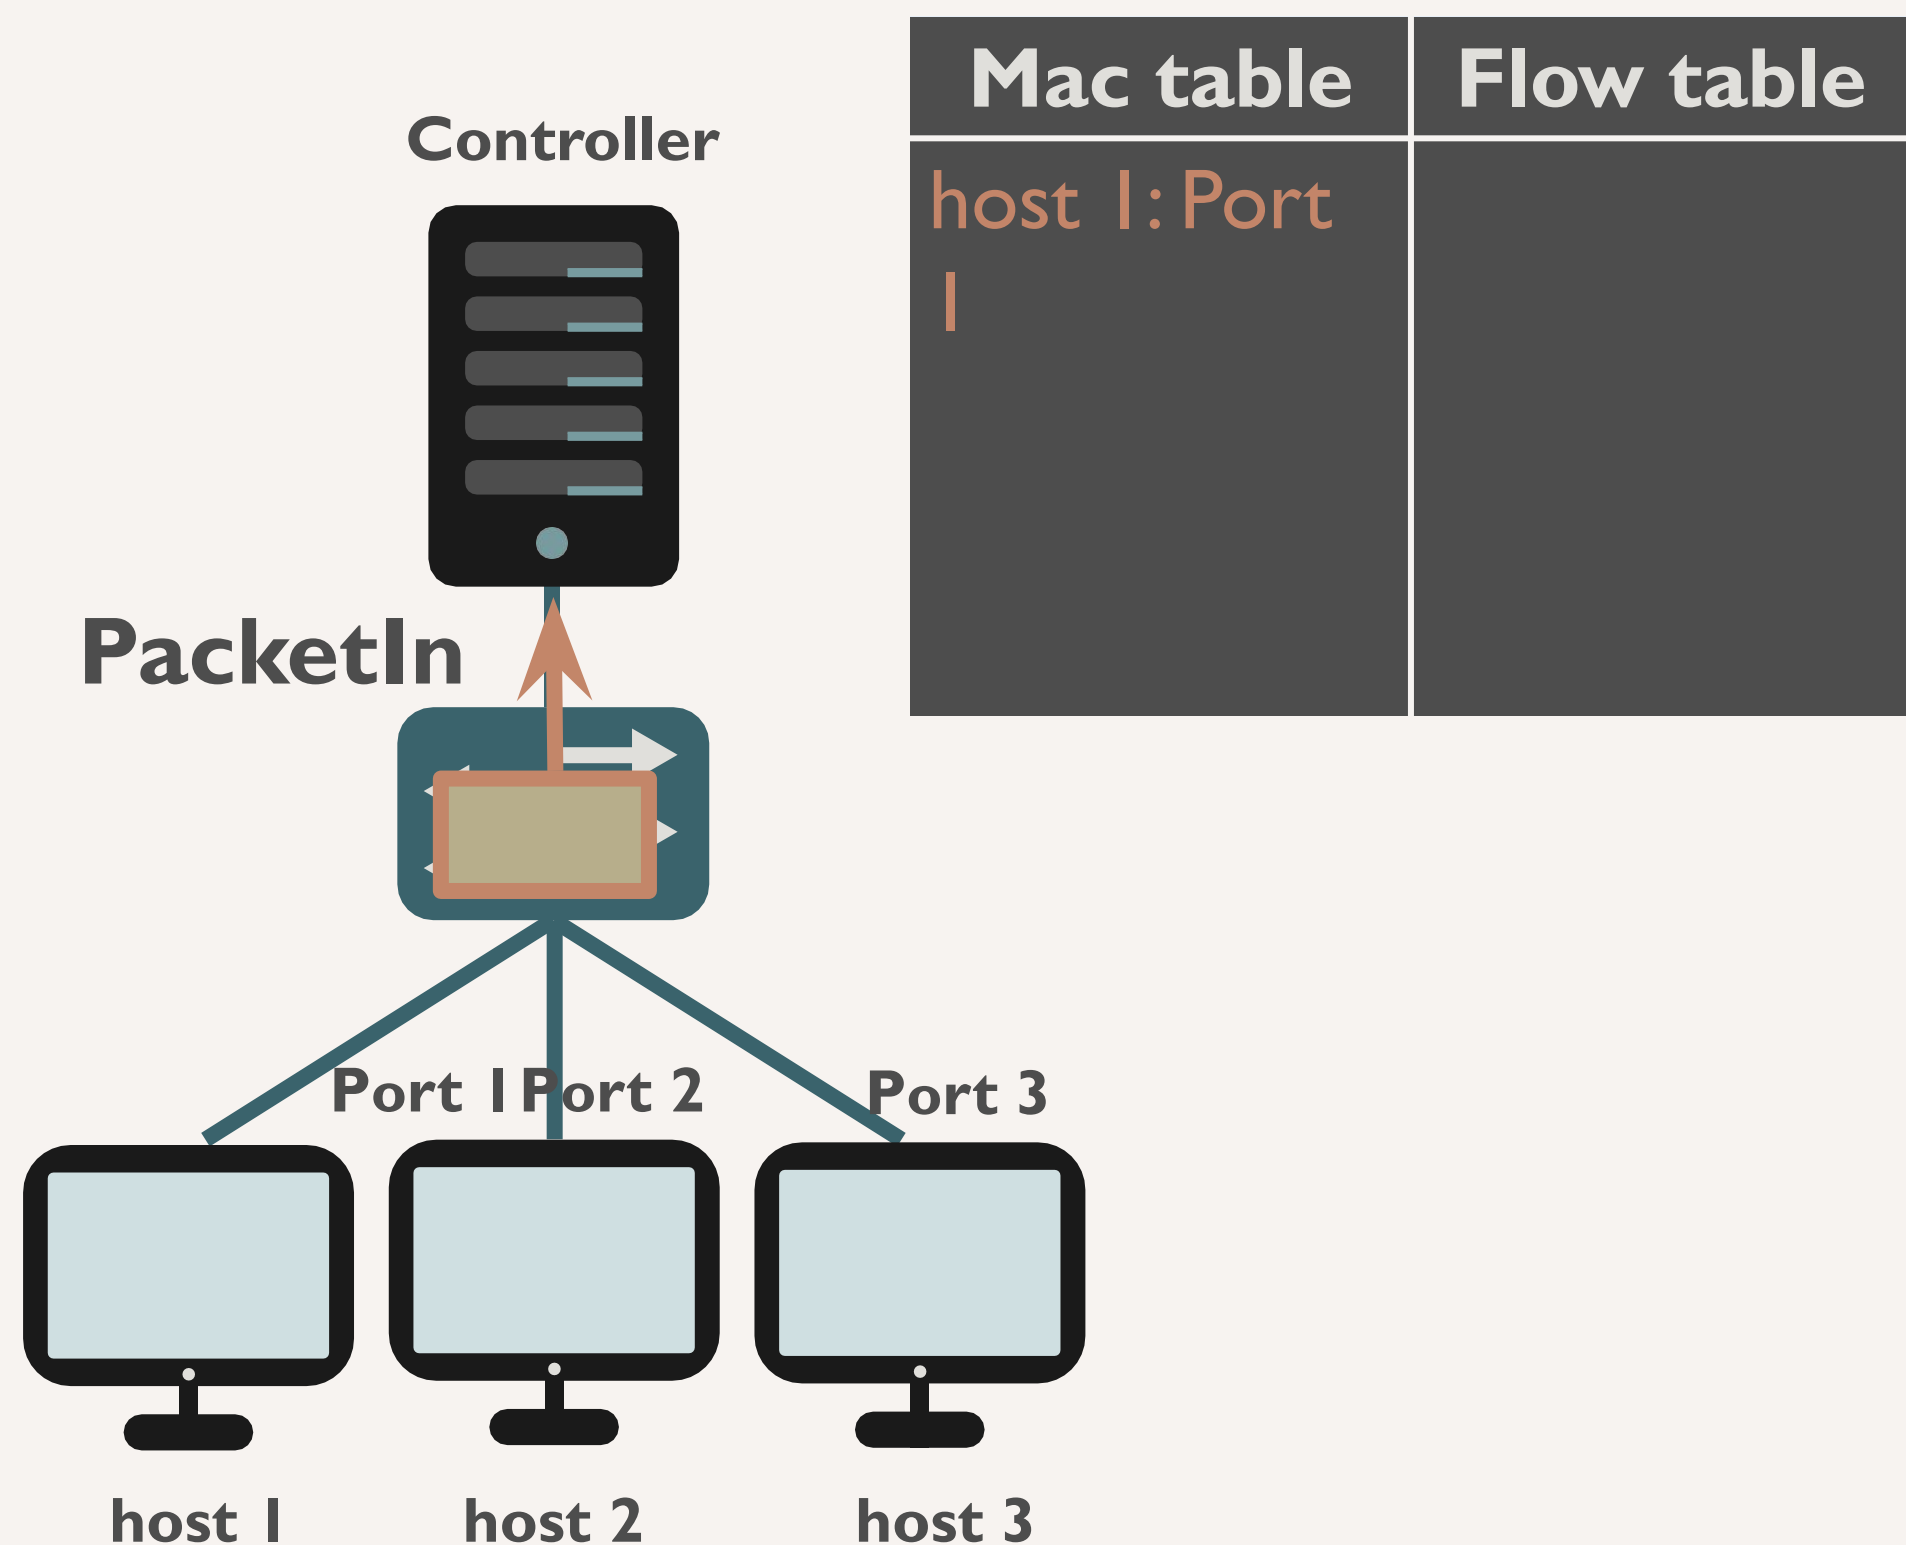
- It will trigger a table miss because the ARP request hasn't been seen by the switch before.



An Example of the Process

– Step 4: PacketIn

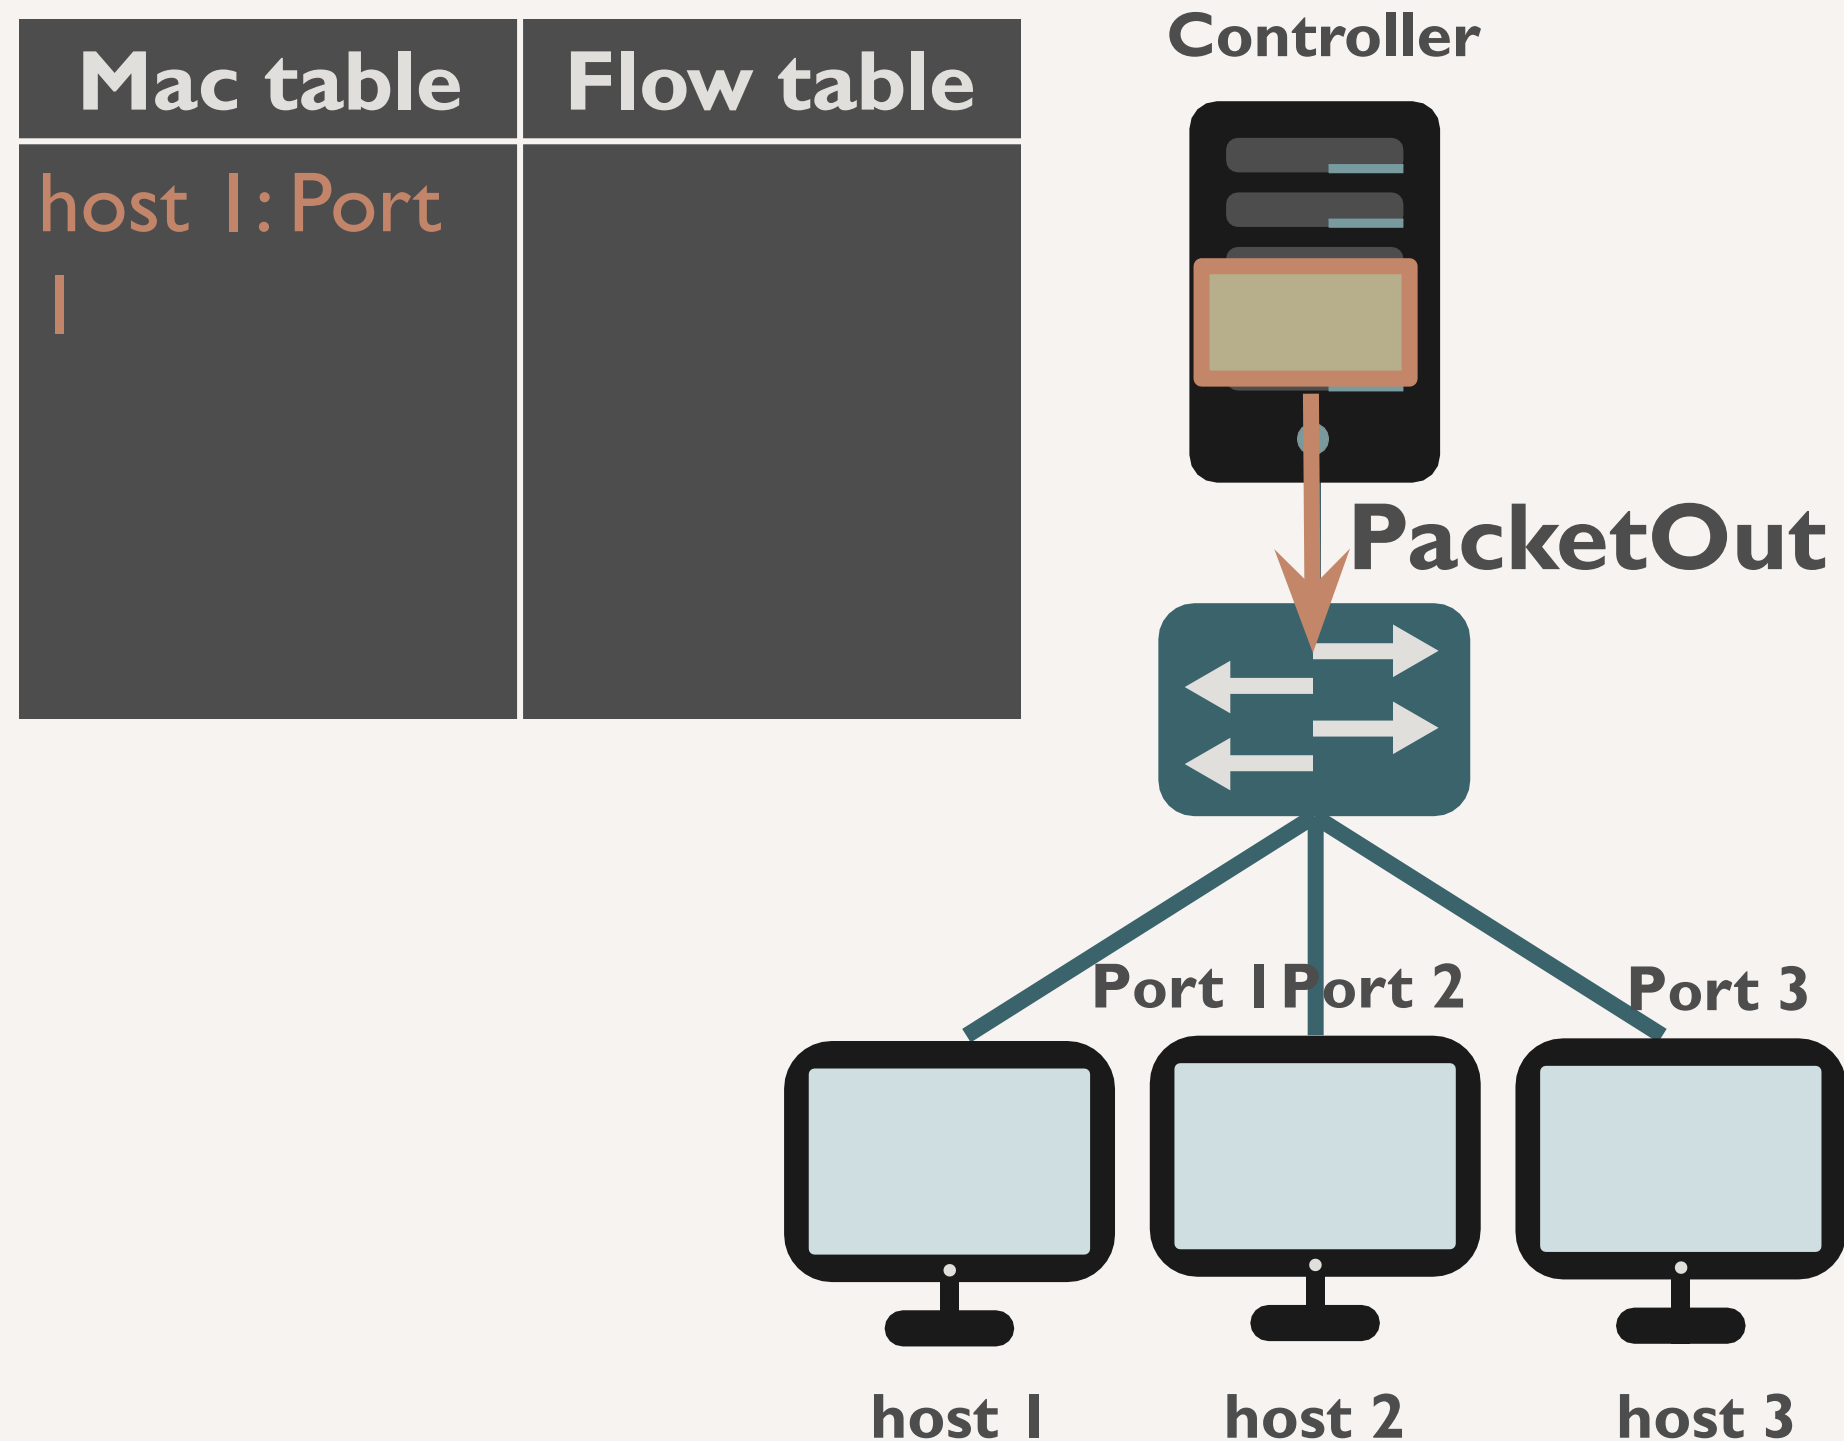
- Thus, the switch will send a **PacketIn** to the controller.



An Example of the Process

– Step 5: PacketOut

- Controller learned the source host and then it **flooded the packet** by **PacketOut** to **broadcast** the ARP request.



An Example of the Process

– Step 5: PacketOut

```
def _packet_in_handler(self, ev):
    #...
    in_port = msg.match['in_port']
    #...
    self.logger.info("packet in %s %s %s %s", dpid, src, dst, in_port)

    # learn a mac address to avoid FLOOD next time.
    self.mac_to_port[dpid][src] = in_port

    if dst in self.mac_to_port[dpid]:
        out_port = self.mac_to_port[dpid][dst]
    else:
        out_port = ofproto.OFPP_FLOOD

    actions = [parser.OFPACTIONOutput(out_port)]

    # install a flow to avoid packet_in next time
    if out_port != ofproto.OFPP_FLOOD:
        match = parser.OFPMATCH(in_port=in_port, eth_dst=dst)
        self.add_flow(datapath, 1, match, actions)
    #...
```

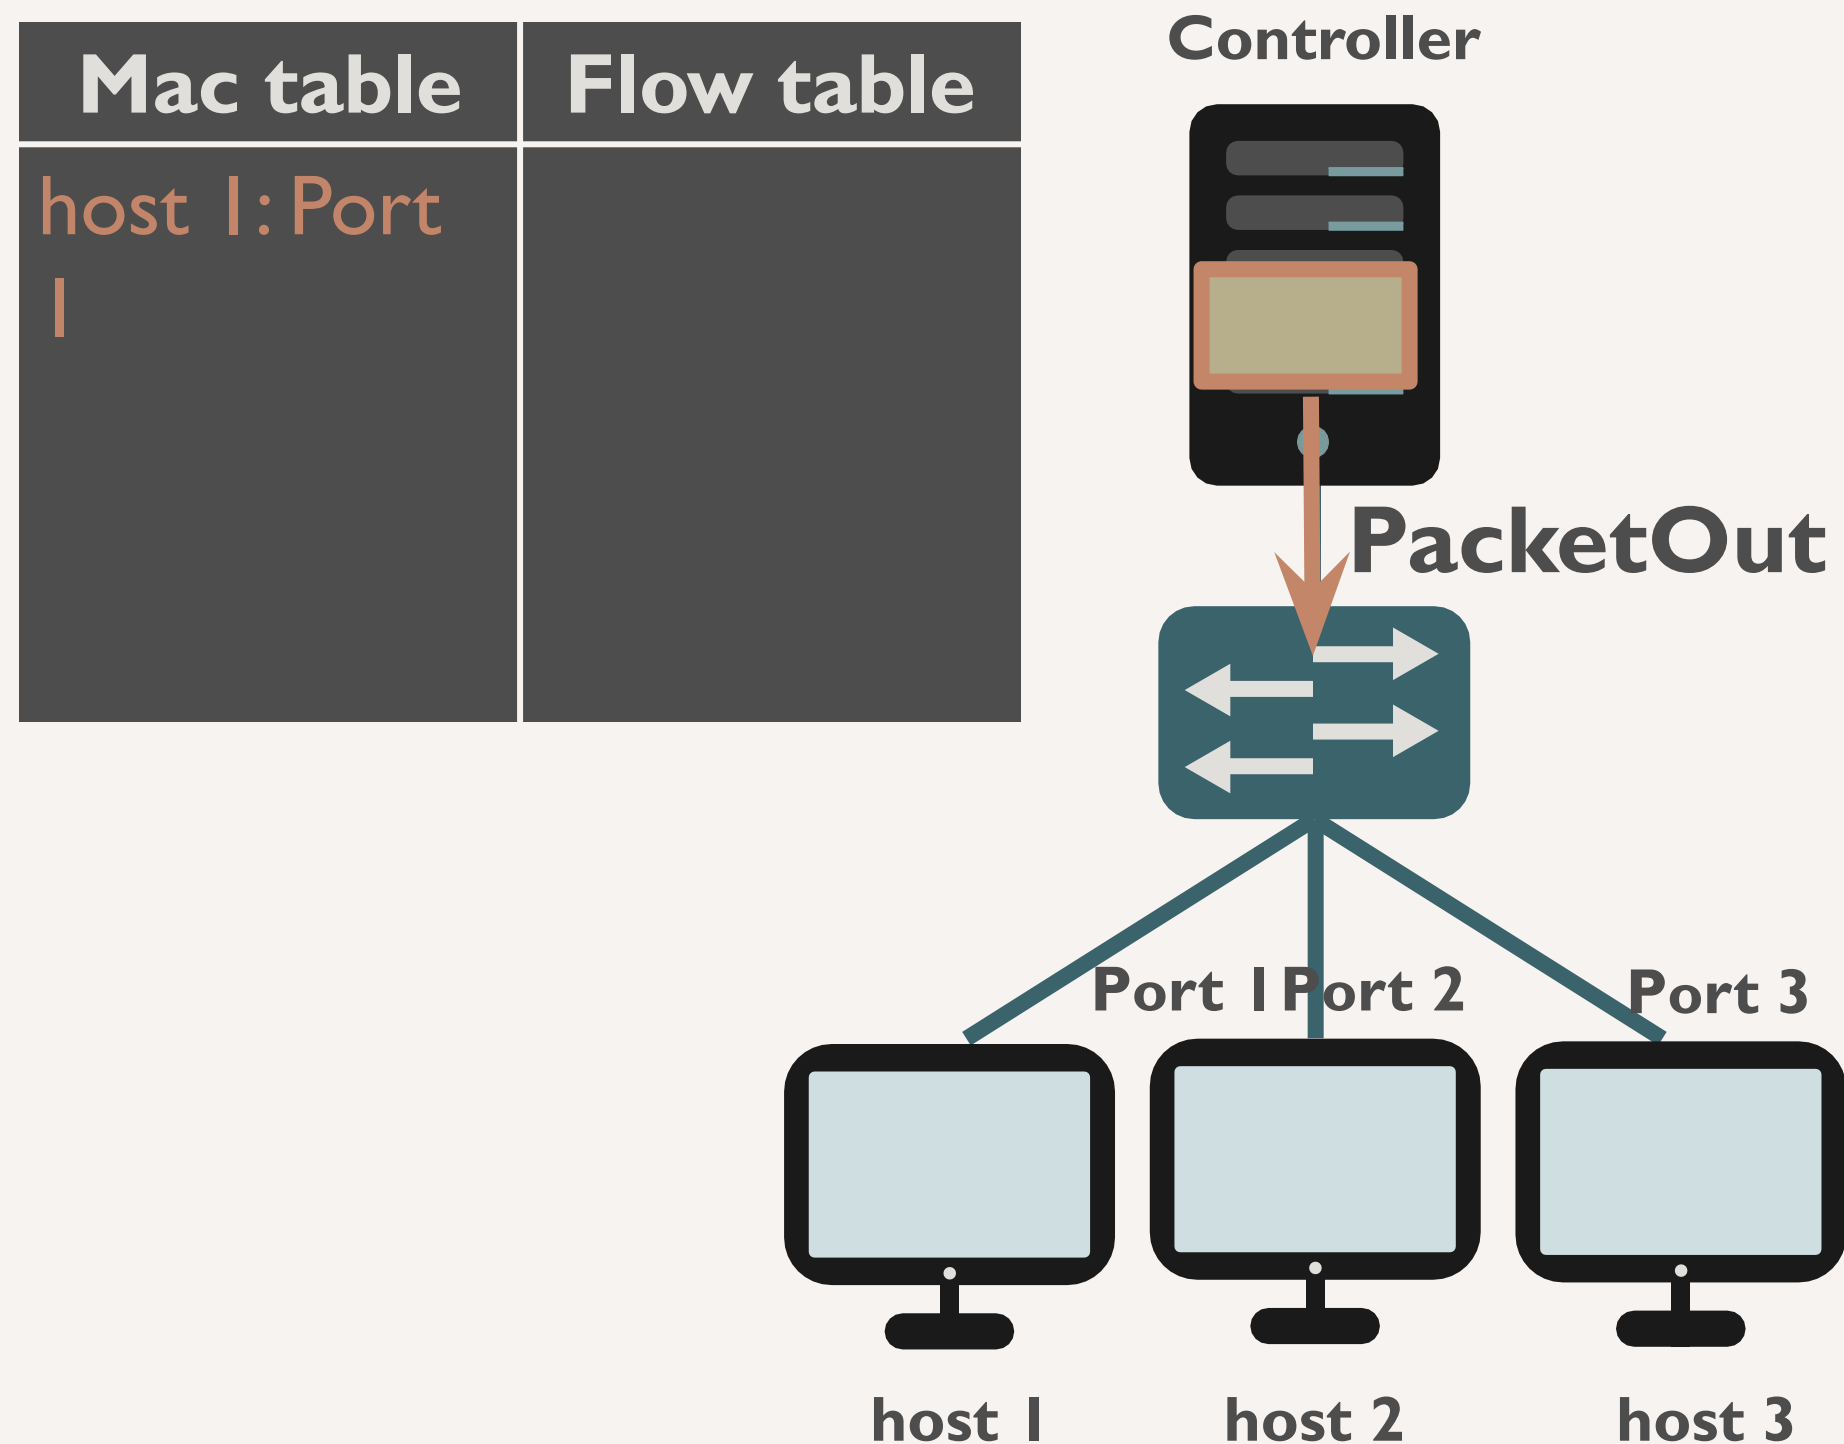
**Learn the
mac addr.
of in_port.**

**Flood the
packet.**

An Example of the Process

– Step 5: PacketOut

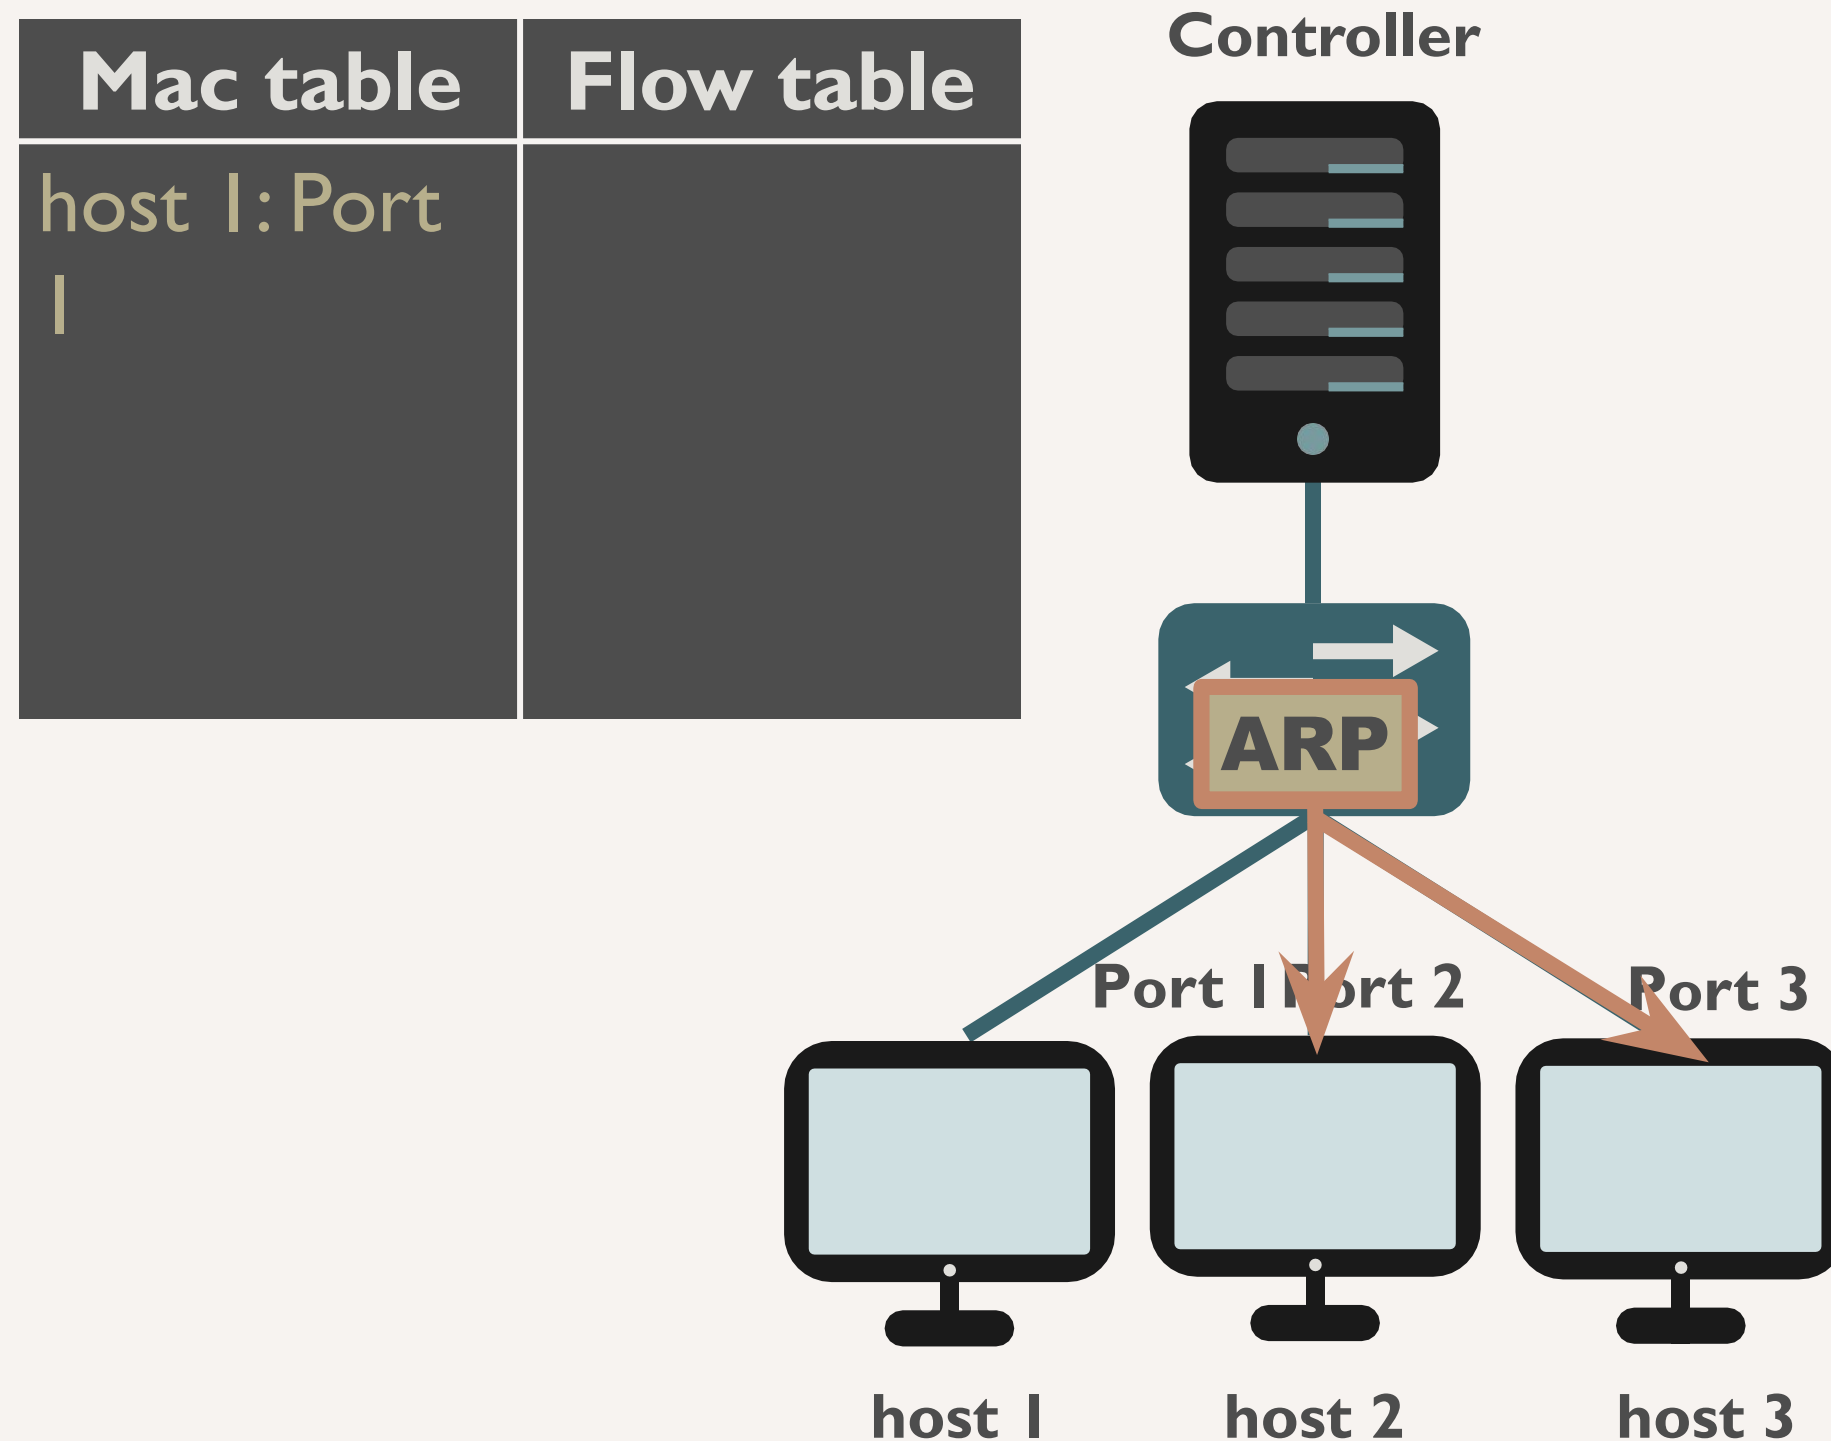
- Controller learned the source host and then it **flooded the packet** by **PacketOut** to **broadcast** the ARP request.



An Example of the Process

– Step 5: PacketOut

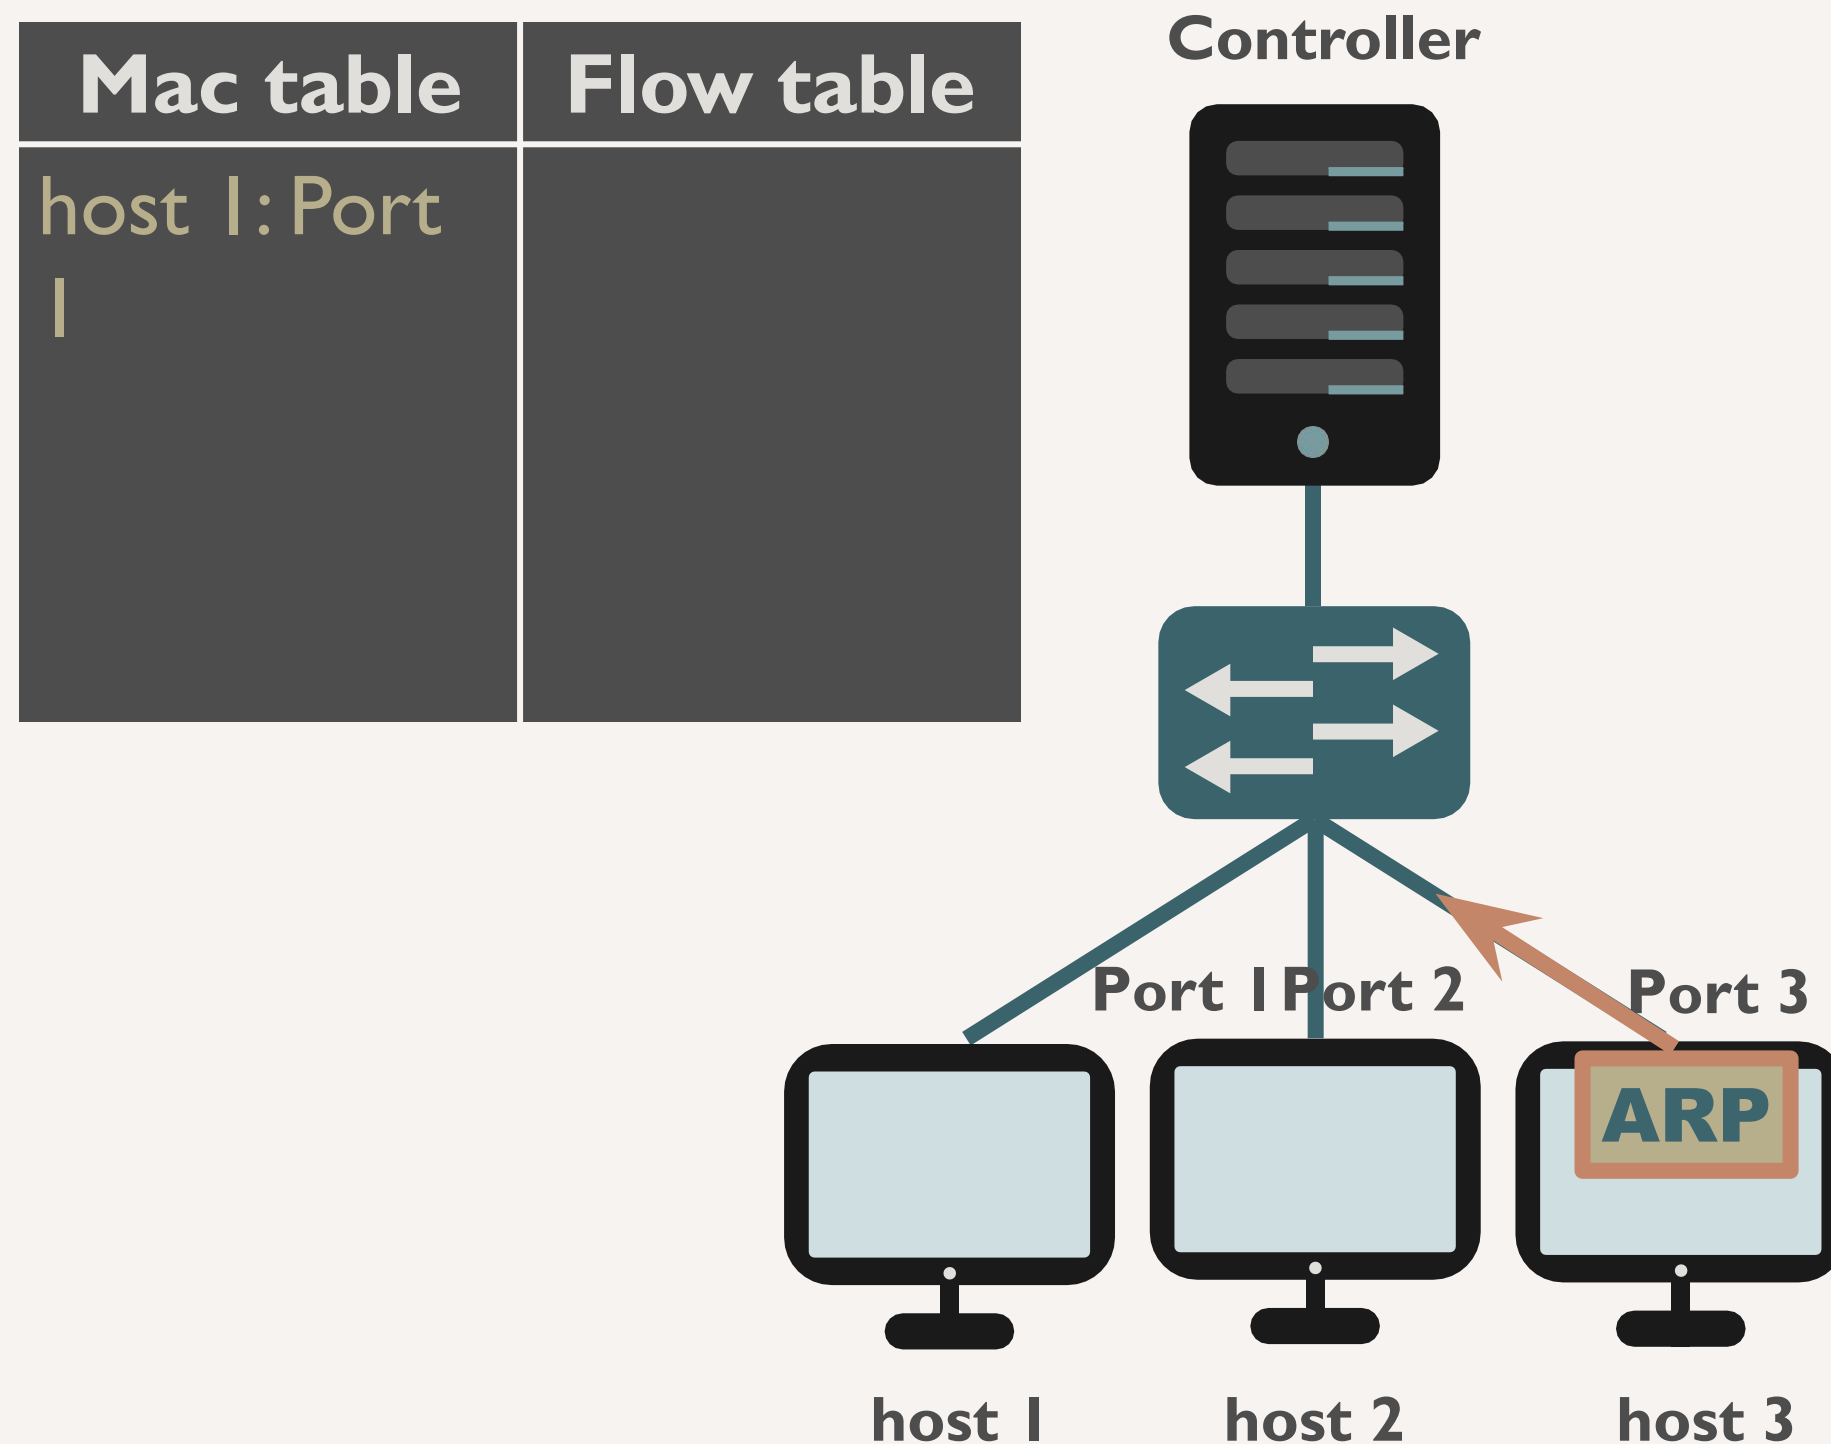
- Controller learned the source host and then it **flooded the packet** by **PacketOut** to **broadcast** the ARP request.



An Example of the Process

– Step 6: PacketIn again

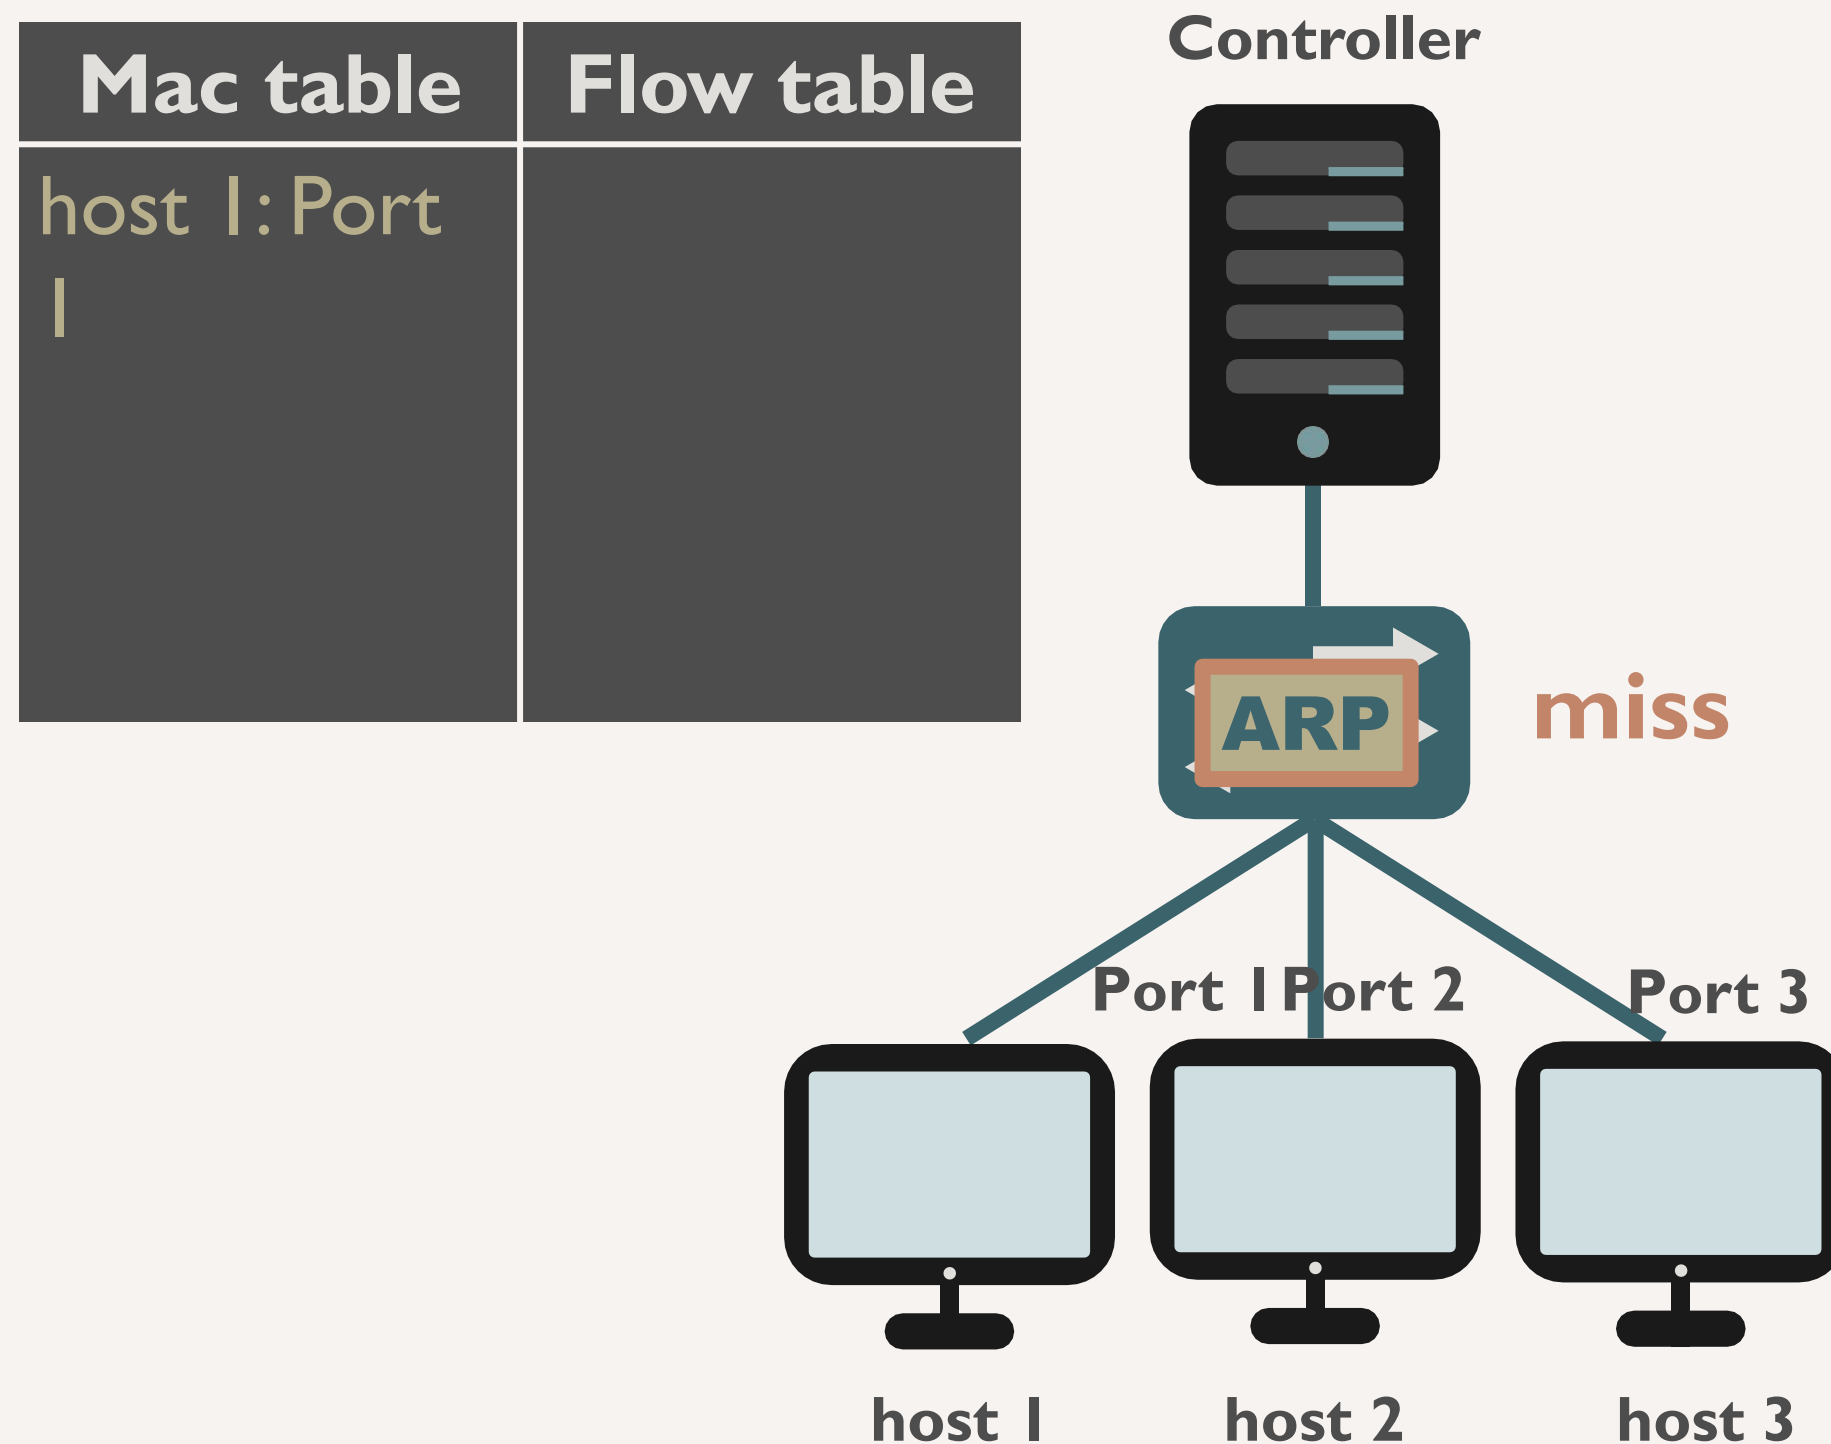
- The destination host (host 3) will reply the ARP request using an **ARP replay** and trigger another **PacketIn**.



An Example of the Process

– Step 6: PacketIn again

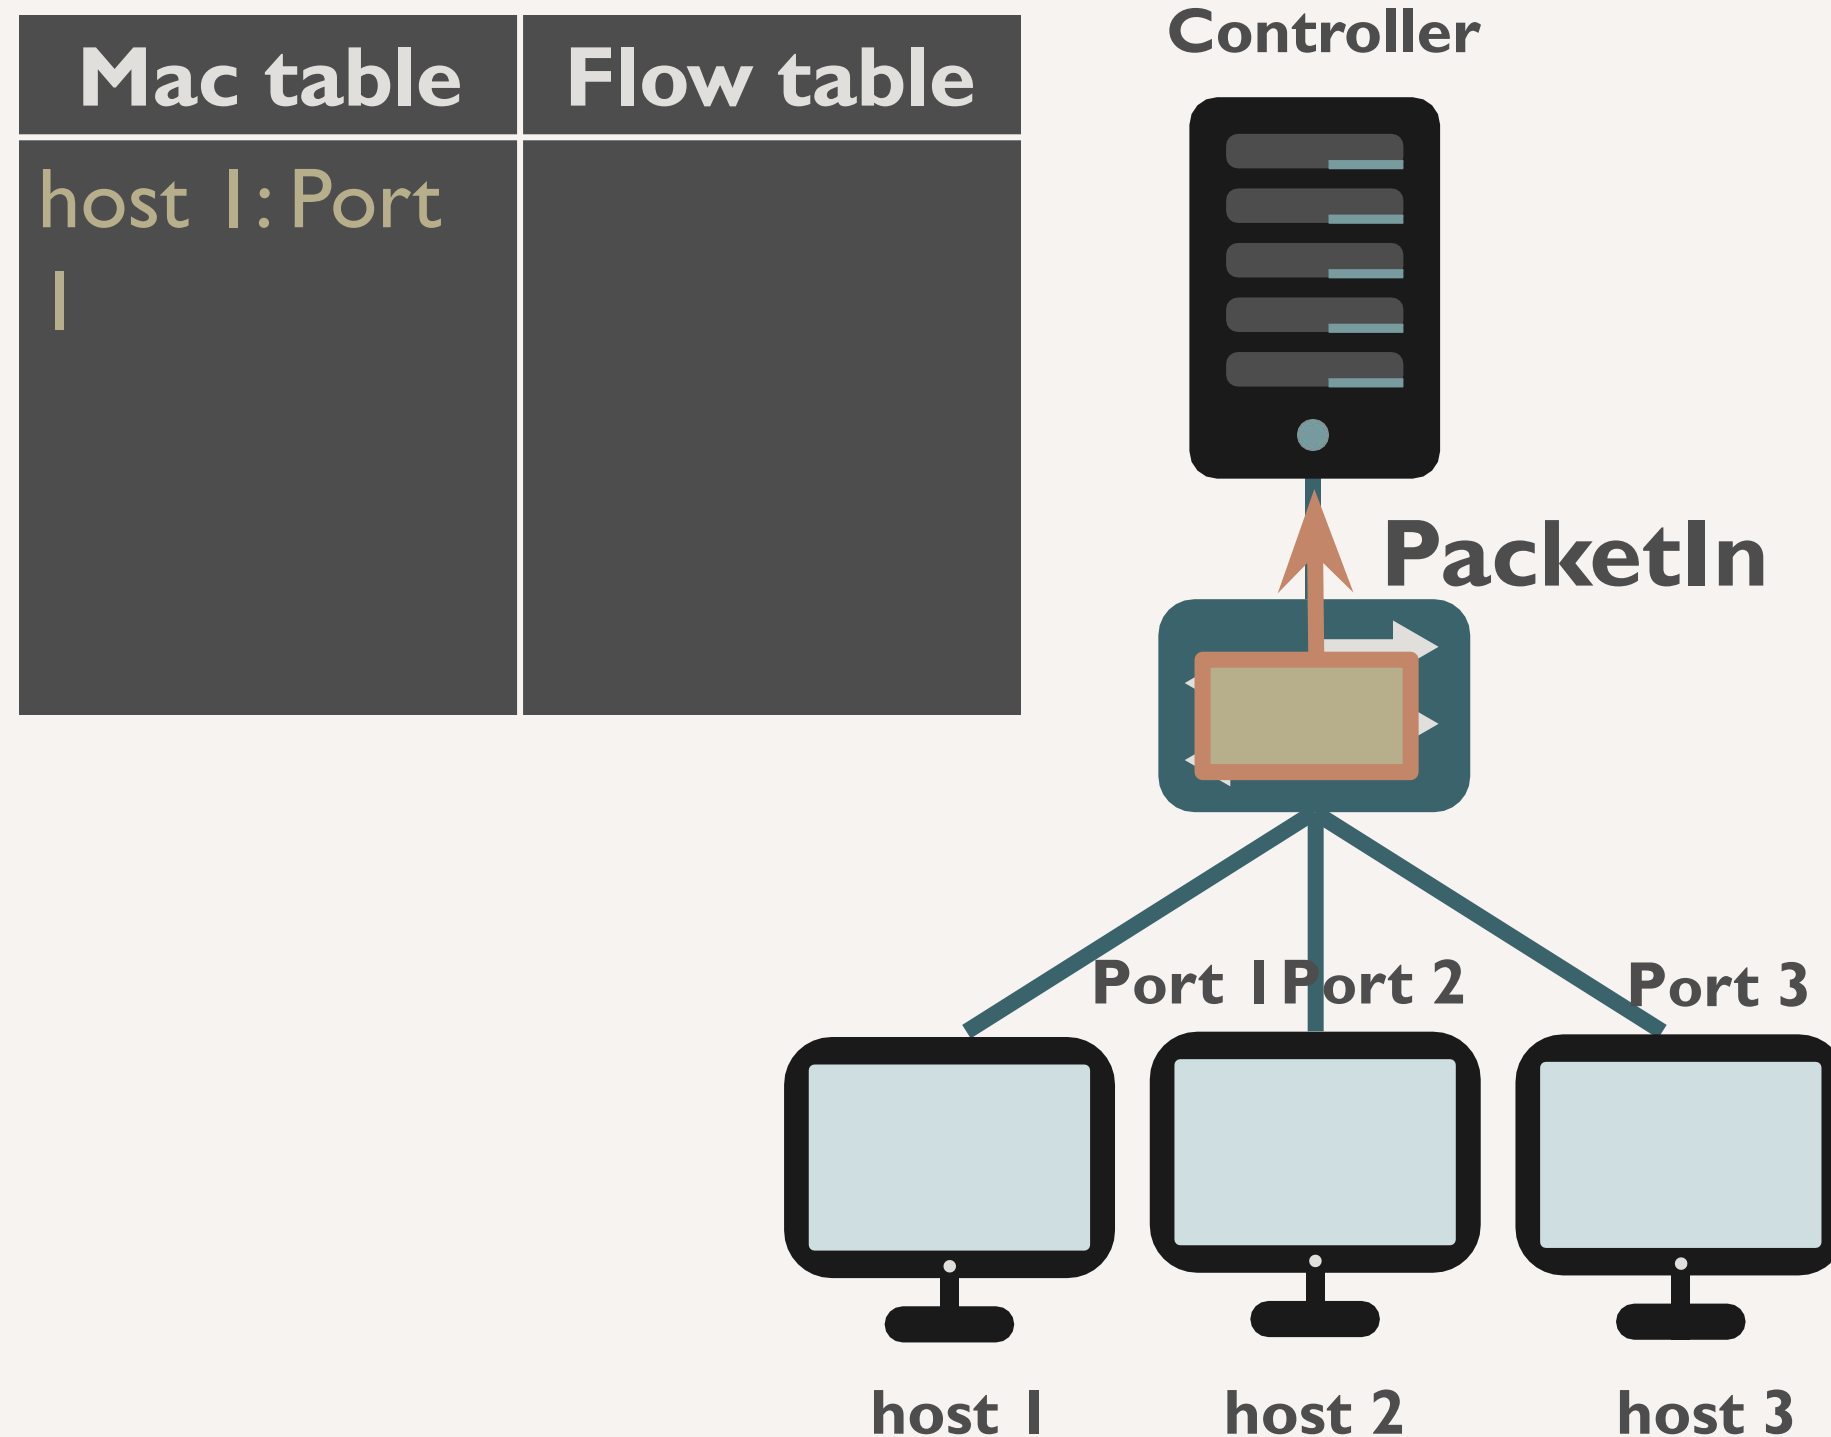
- The destination host (host 3) will reply the ARP request using an **ARP replay** and trigger another **PacketIn**.



An Example of the Process

– Step 6: PacketIn again

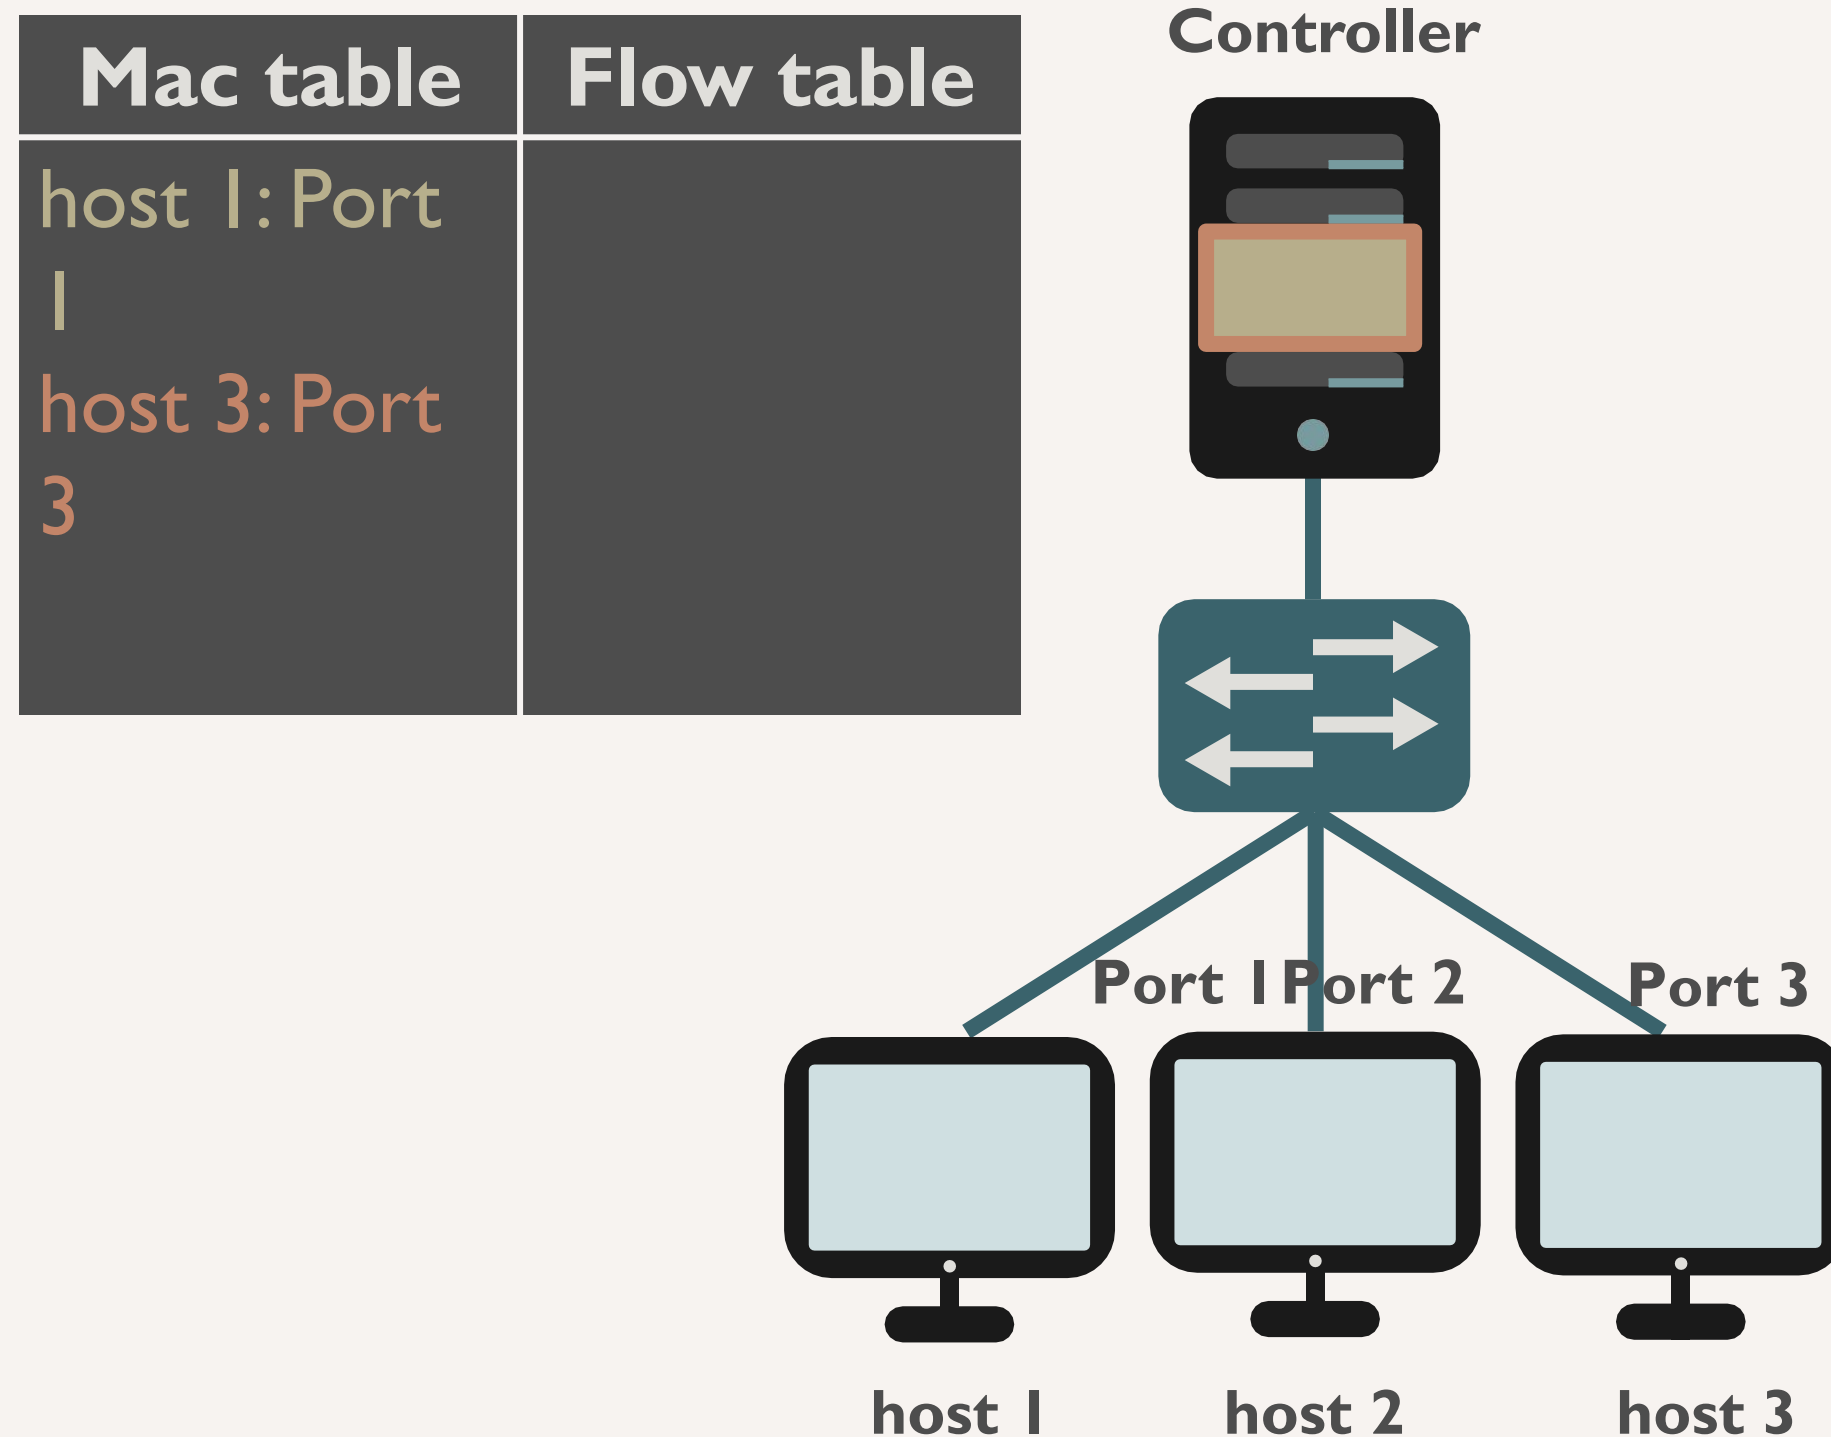
- The destination host (host 3) will reply the ARP request using an **ARP replay** and trigger another **PacketIn**.



An Example of the Process

– Step 6: PacketIn again

- The destination host (host 3) will reply the ARP request using an **ARP replay** and trigger another **PacketIn**.



An Example of the Process

– Step 6: PacketIn again

- In `_packet_in_handler`

```
# install a flow to avoid packet_in nexttime
if out_port != ofproto.OFPP_FLOOD:
    match = parser.OFPMatch(in_port=in_port, eth_dst=dst)
    self.add_flow(datapath, 1, match, actions)
#...
```

- In `add_flow`

```
def add_flow(self, datapath, priority, match, actions):
    ofproto = datapath.ofproto
    parser = datapath.ofproto_parser
    inst = [parser.OFPIInstructionActions(ofproto.OFPIT_APPLY_ACTIONS,
                                         actions)]

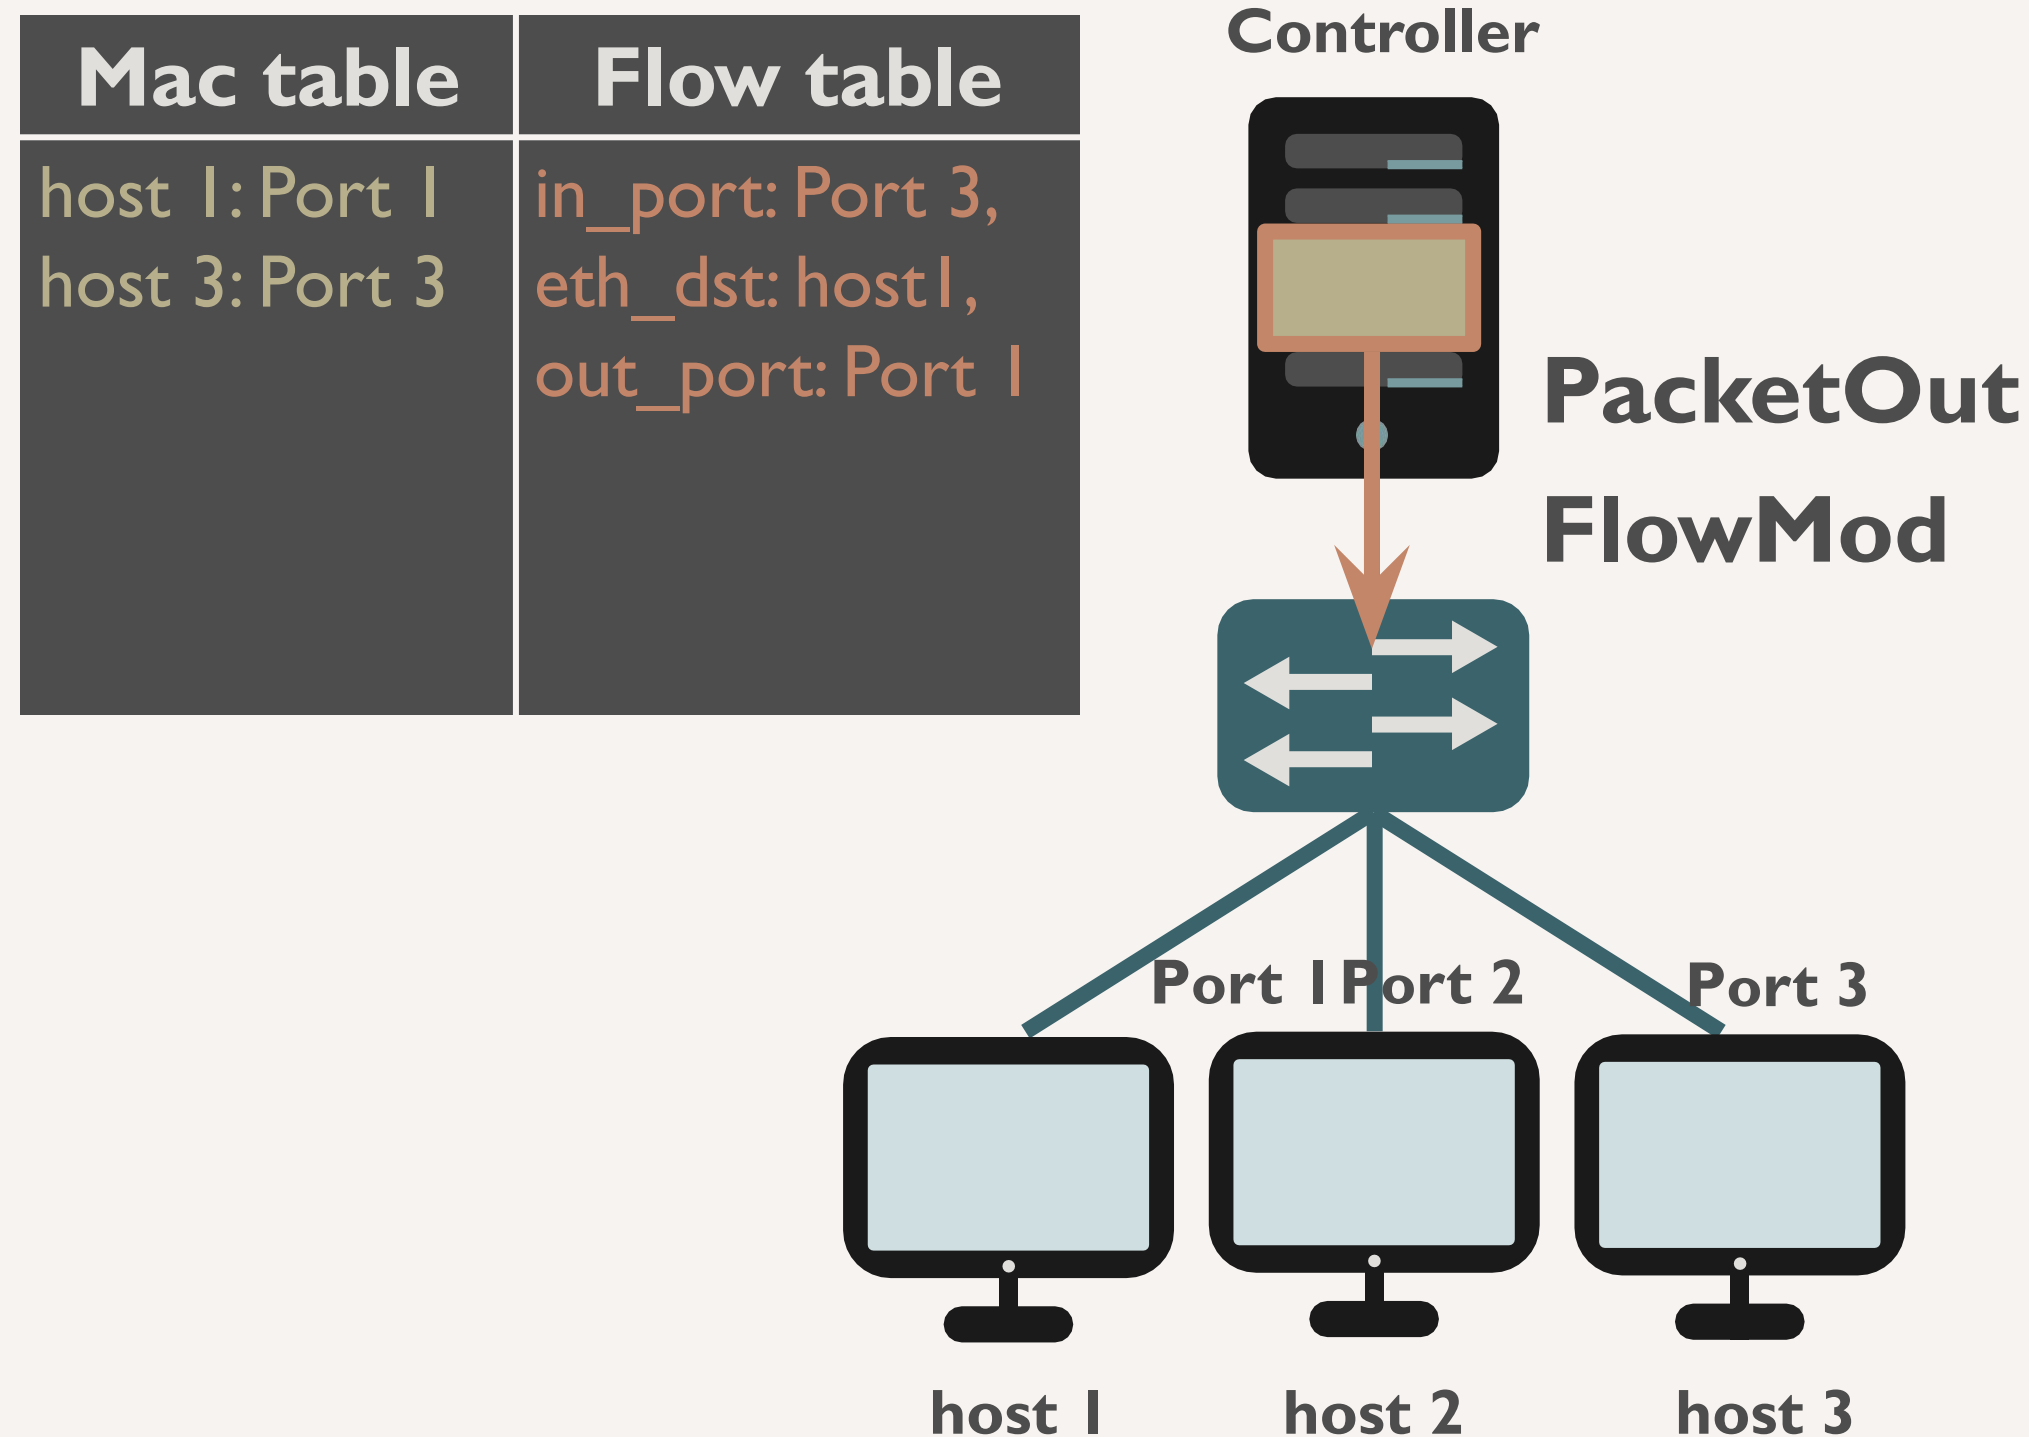
    mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
                             match=match, instructions=inst)
    datapath.send_msg(mod)
```

datapath: which switch
priority: the higher value, the higher priority
match: the rule set that want to match
actions: the action set

An Example of the Process

– Step 7: FlowMod & PacketOut

- This time, the controller knows the destination and source. So it adds a rule using **flow mod** along with the **Packet out**.




An Example of the Process

– Step 7: FlowMod & PacketOut

- This time, the controller knows the destination and source. So it adds a rule using **flow mod** along with the **Packet out**.

Mac table	Flow table
host 1: Port 1 host 3: Port 3	in_port: Port 3, eth_dst: host 1, out_port: Port 1

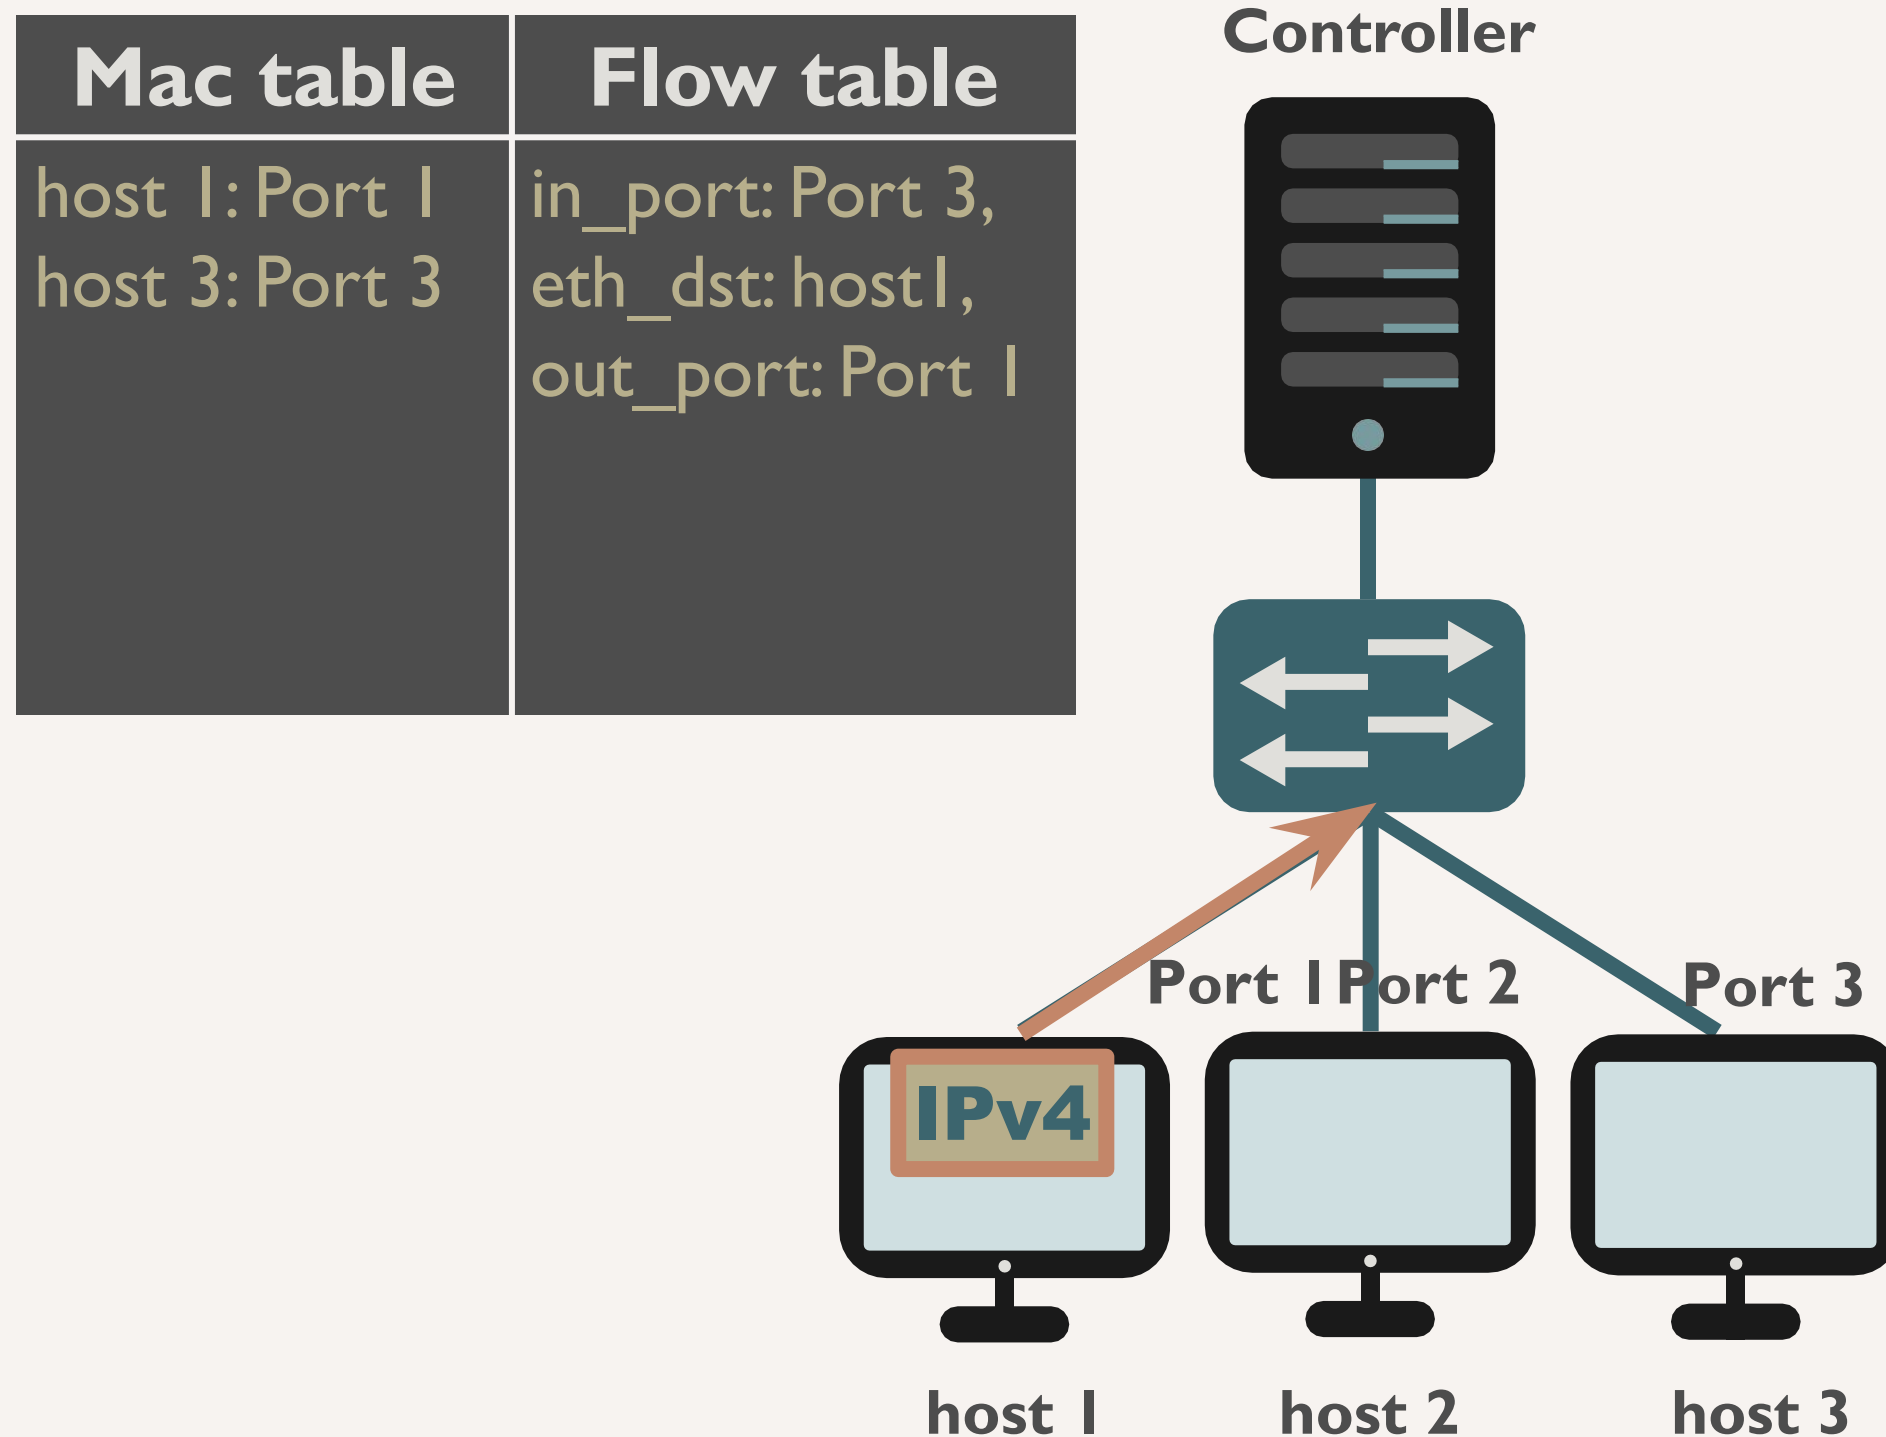
Learned the MAC



An Example of the Process

– Step 8: PacketIn of the 2nd iteration

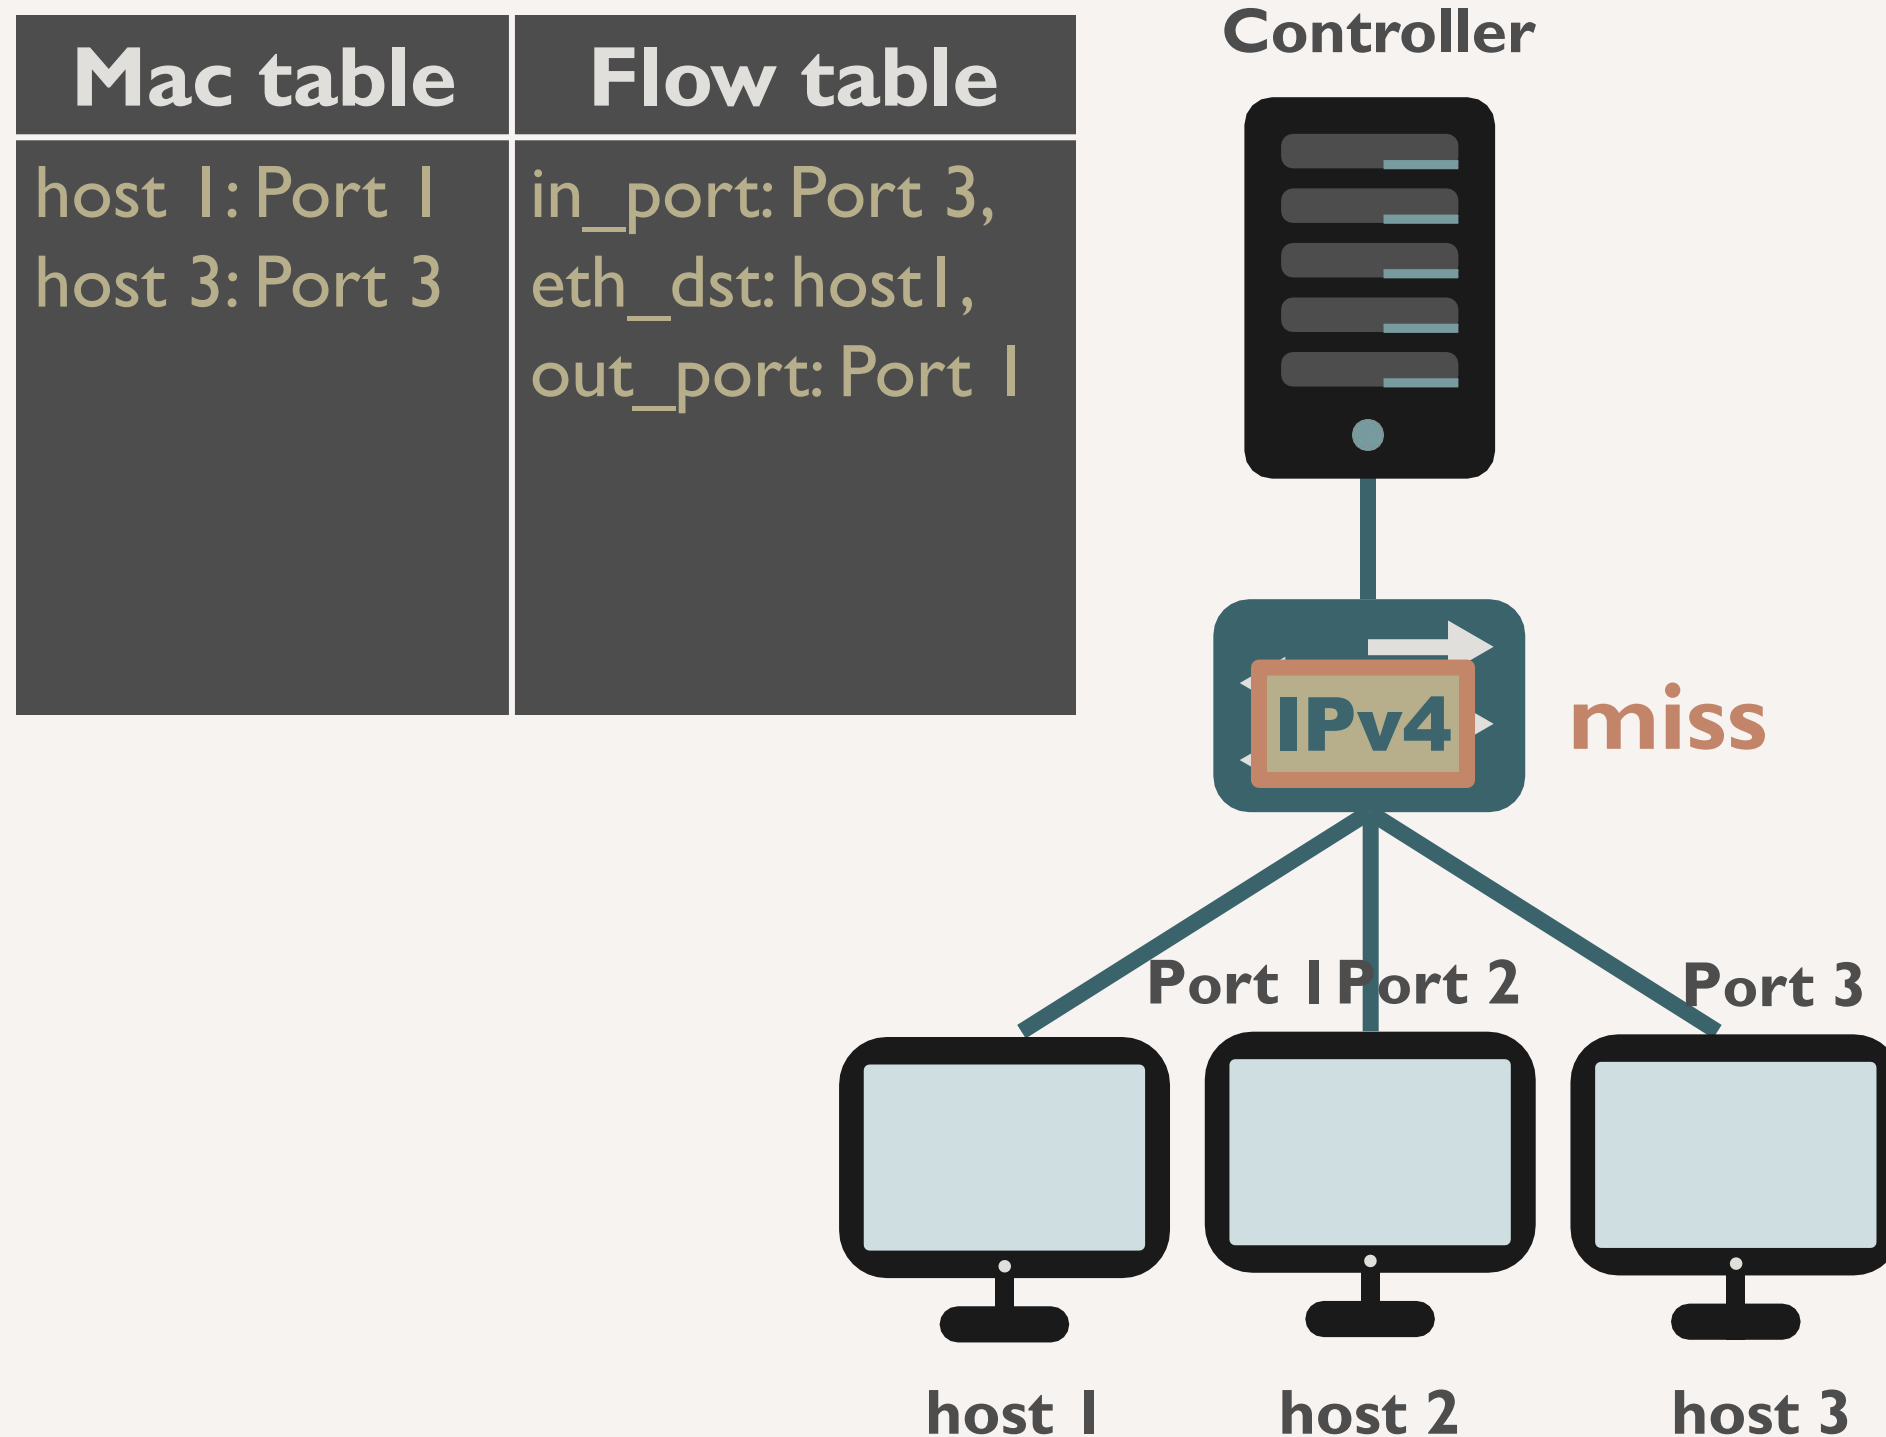
- The next ping for host 1 to host 3 inside the network.



An Example of the Process

– Step 8: PacketIn of the 2nd iteration

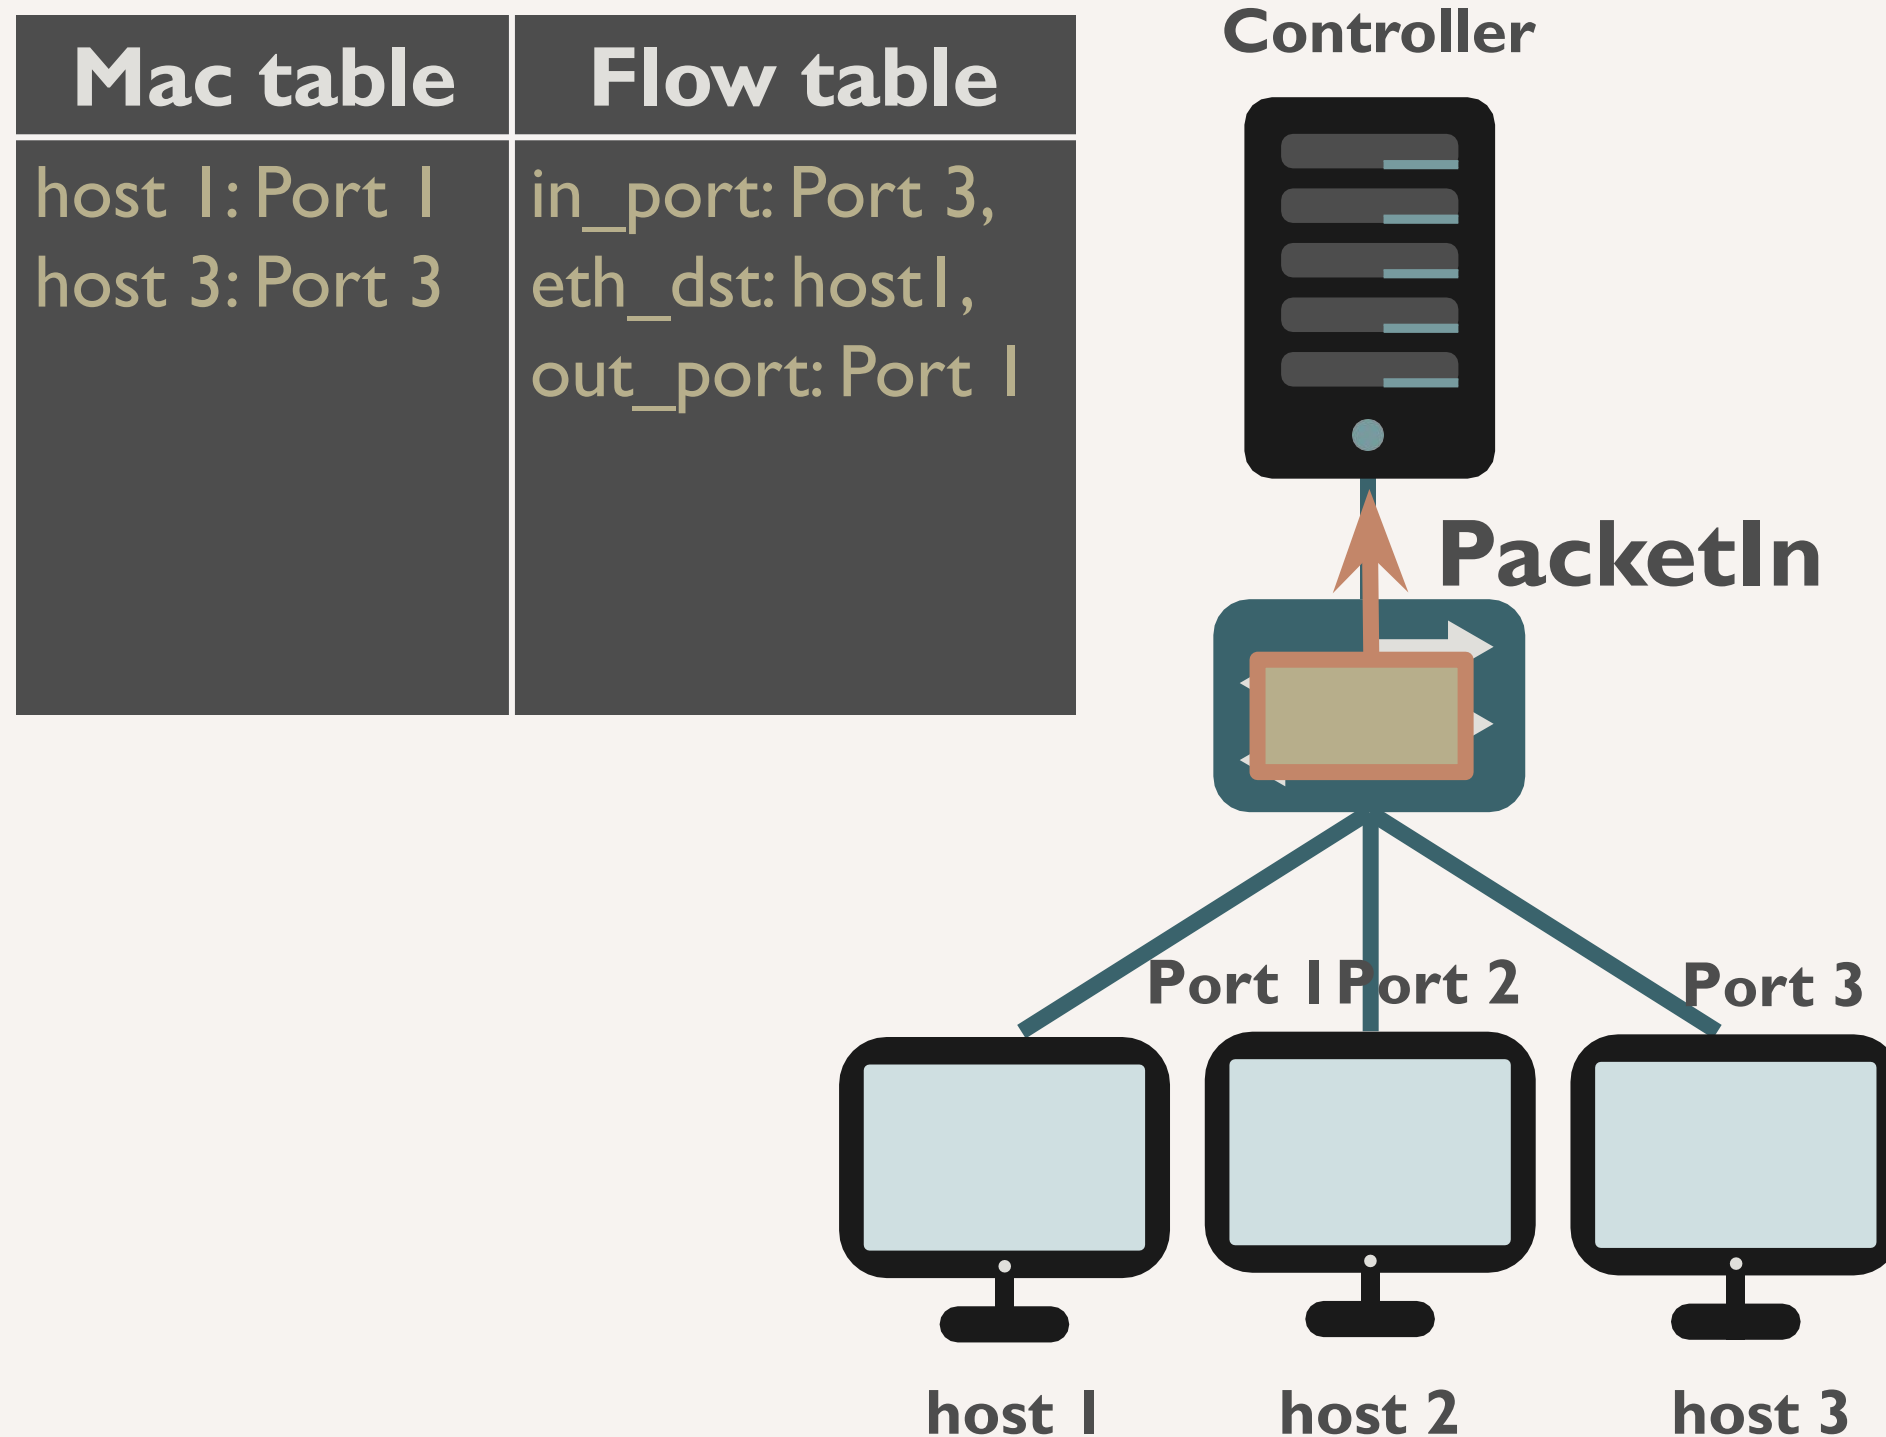
- It will trigger a table miss because it's the first time for the switch to get a packet sent from host 1 to host 3.



An Example of the Process

– Step 8: PacketIn of the 2nd iteration

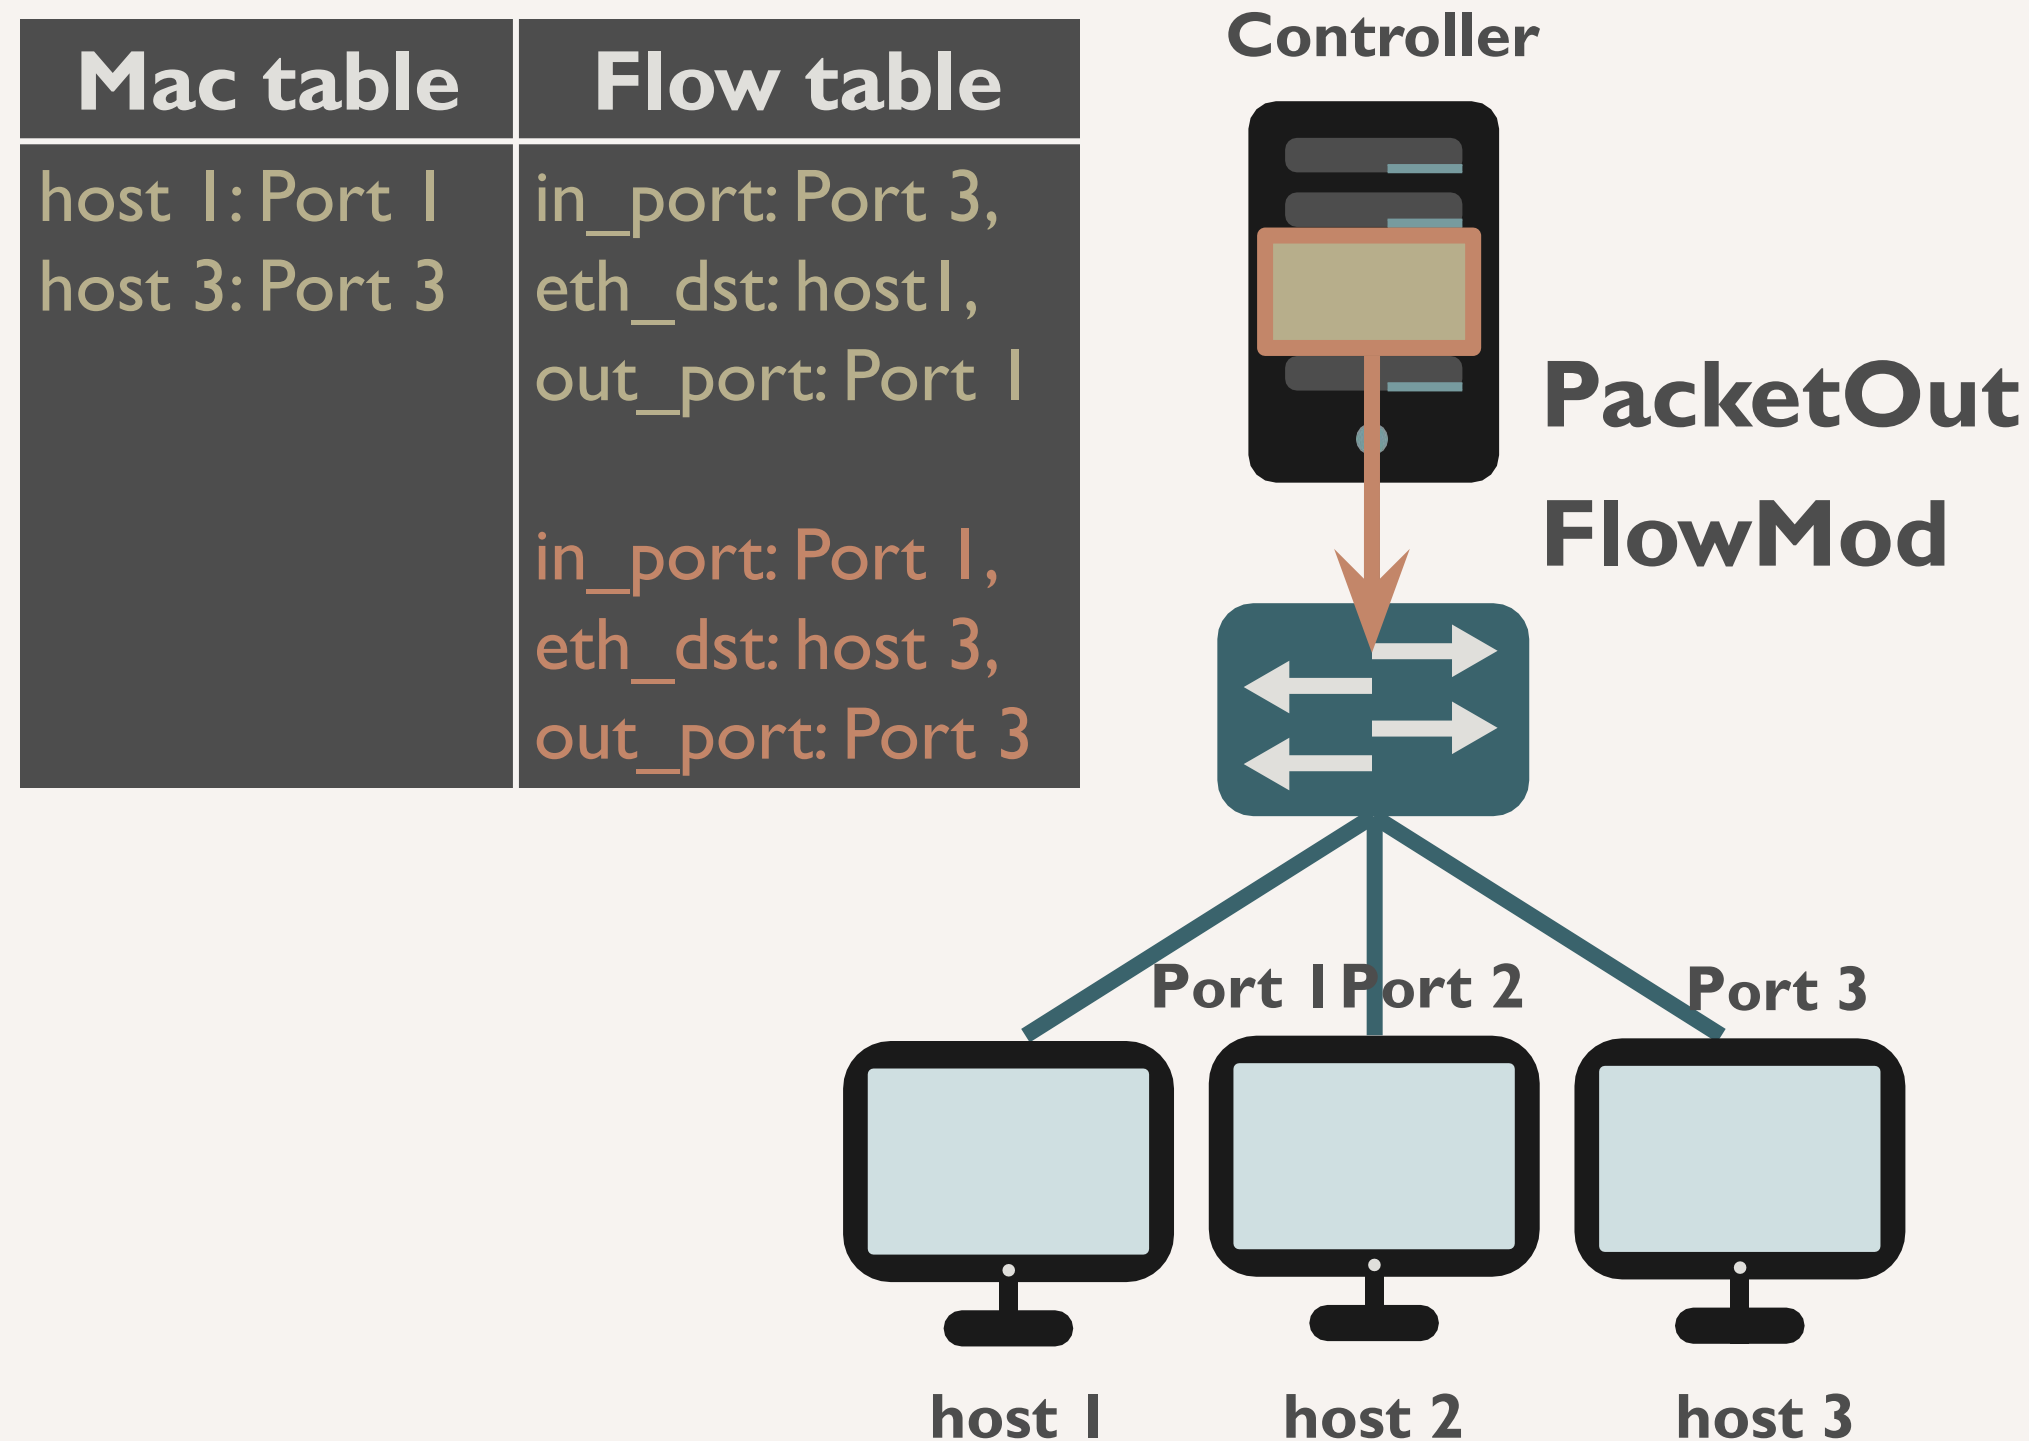
- It will trigger a table miss because it's the first time for the switch to get a packet sent from host 1 to host 3.



An Example of the Process

– Step 8: PacketIn of the 2nd iteration

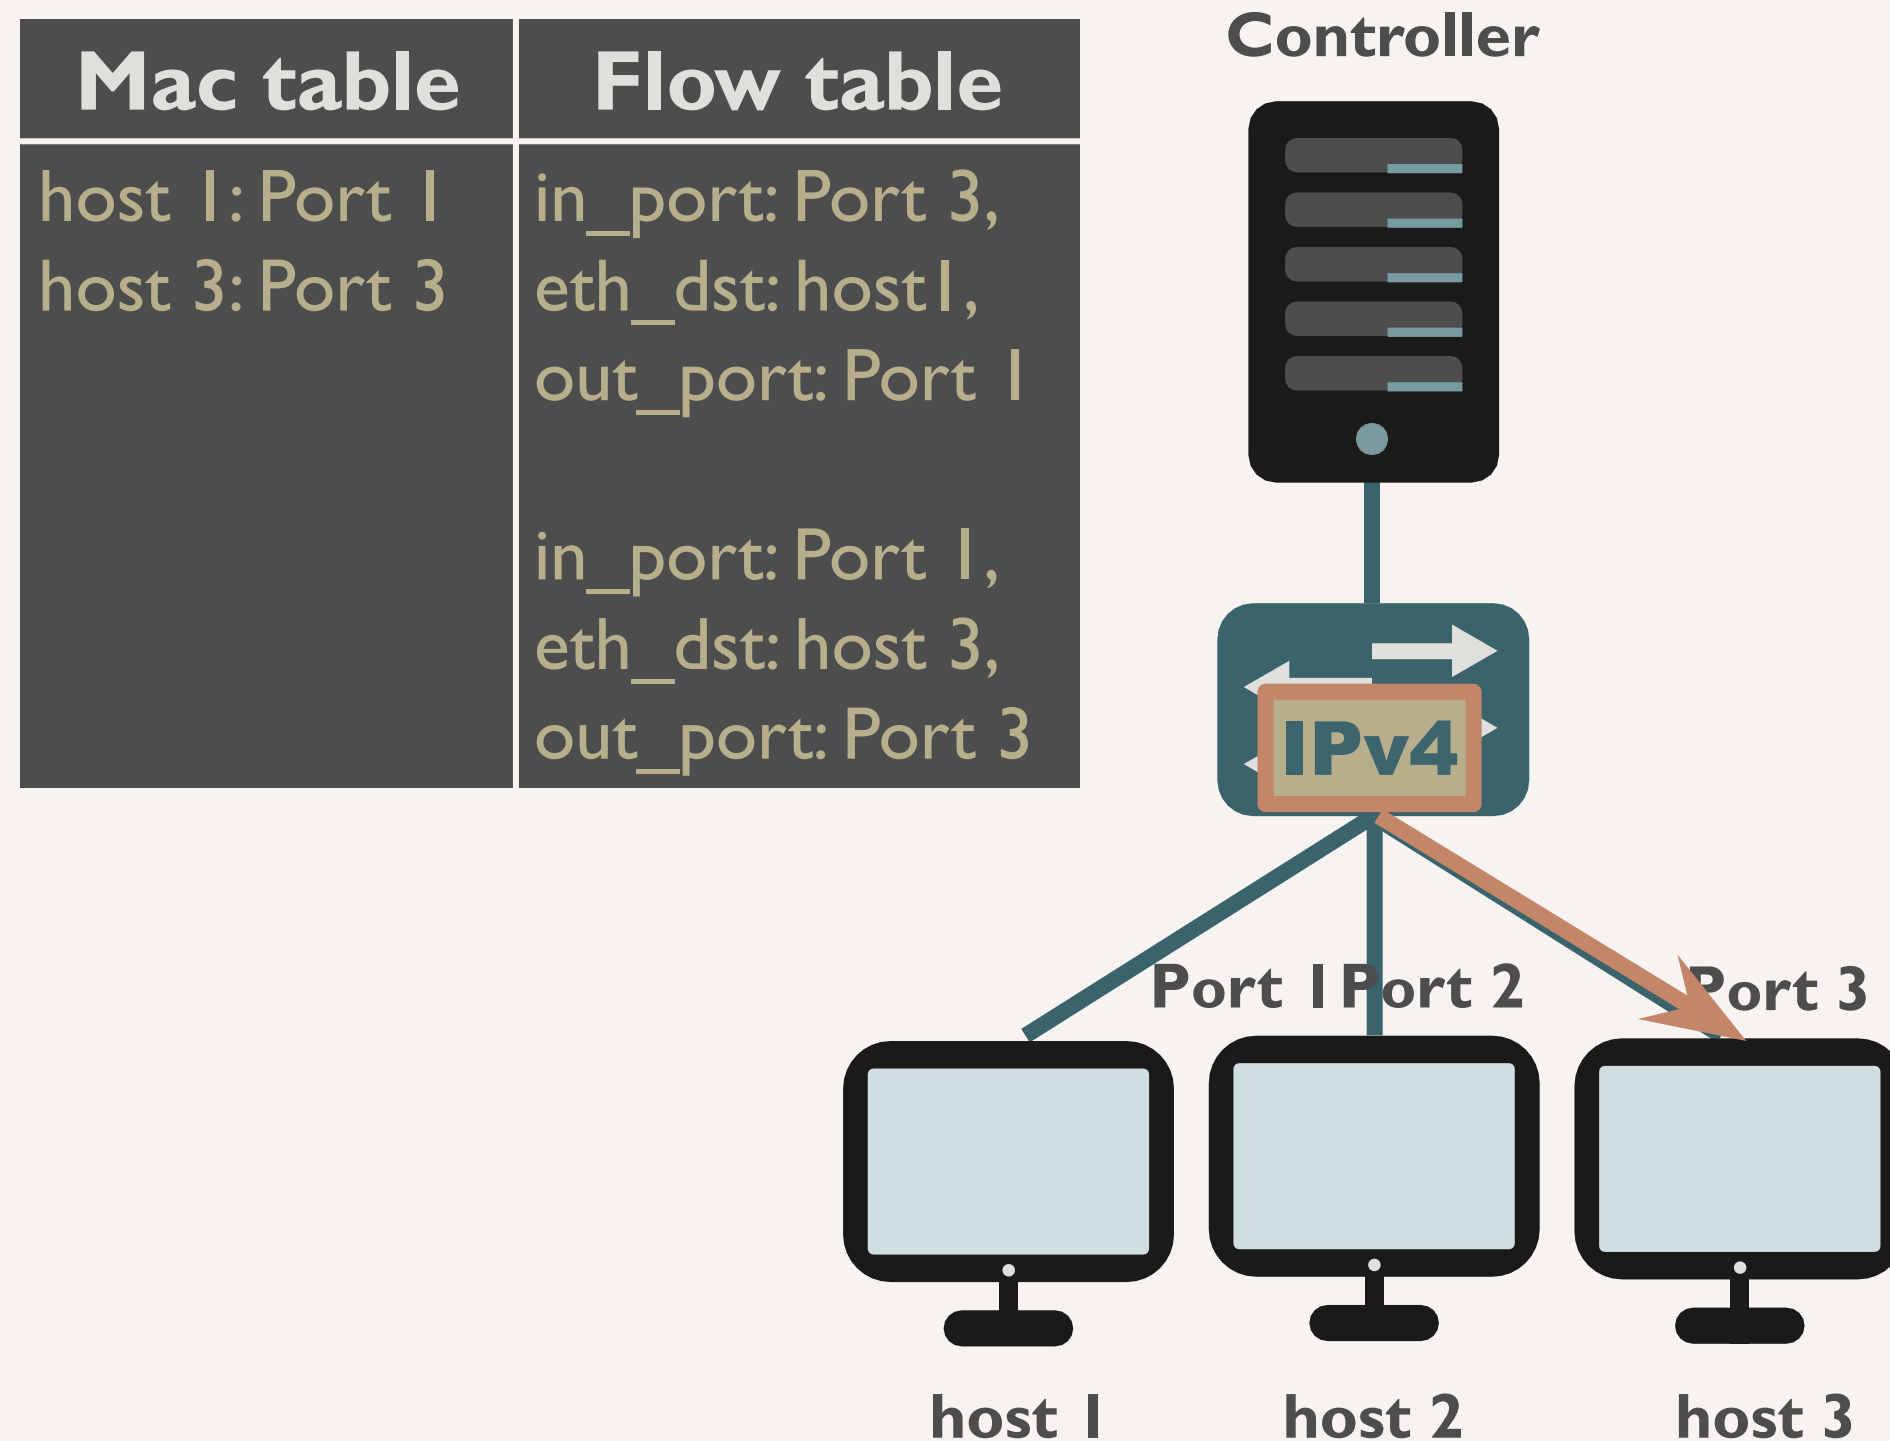
- This time, the controller knows the destination and the source. So it **adds a rule using FlowMod** and send a **PacketOut**.



An Example of the Process

– Step 8: PacketIn of the 2nd iteration

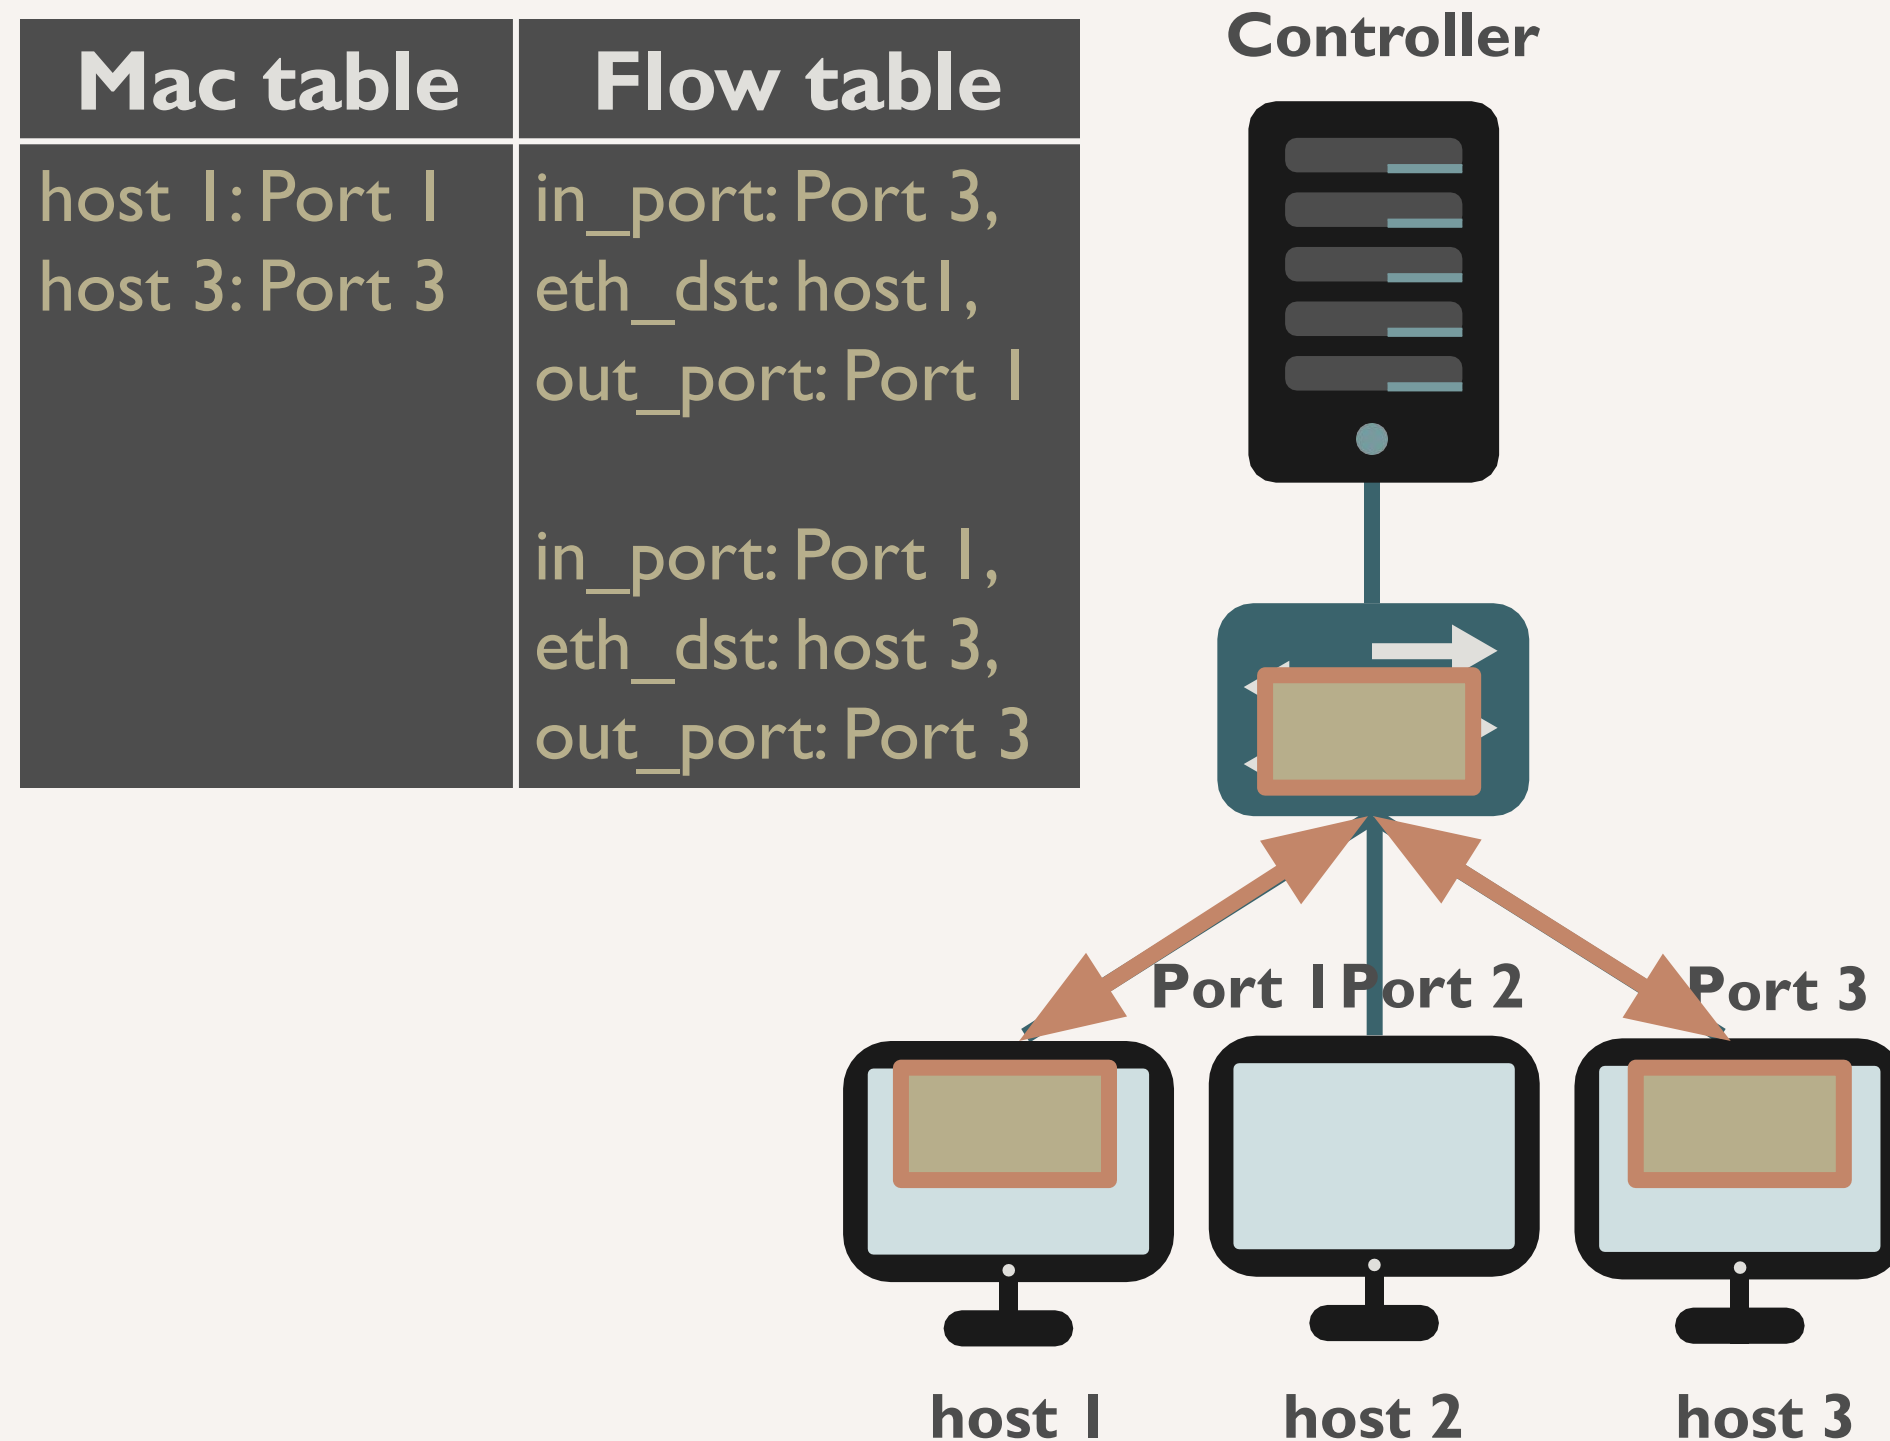
- This time, the controller knows the destination and the source. So it **adds a rule using FlowMod** and send a **PacketOut**.



An Example of the Process

– Result: the Path Has Made

- The destination host generates reply packet and straightly forward to the source by the switch.



Result - Mininet

```
$ mininet> pingall
```

```
*** Stopping 1 switches
s1
*** Stopping 3 hosts
h1 h2 h3
*** Done
completed in 27.992 seconds
mininet@mininet-vm:~$ ryu-manager --verbose ryu.app.simple_switch_13
ryu-manager: command not found
mininet@mininet-vm:~$ sudo mn --topo single,3 --mac --switch ovsk,protocols=OpenFlow13 --controller
remote
*** Creating network
*** Adding controller
Unable to contact the remote controller at 127.0.0.1:6653
Unable to contact the remote controller at 127.0.0.1:6633
Setting remote controller to 127.0.0.1:6653
*** Adding hosts:
h1 h2 h3
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1) (h3, s1)
*** Configuring hosts
h1 h2 h3
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> ryu-manager --verbose ryu.app.simple_switch_13
*** Unknown command: ryu-manager --verbose ryu.app.simple_switch_13
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3
h2 -> h1 h3
h3 -> h1 h2
*** Results: 0% dropped (6/6 received)
mininet>
```

Result - Ryu

- Format: switch_id, src_mac, dst_mac, inport

```
BRICK ofp_event
  PROVIDES EventOFPSwitchFeatures TO {'SimpleSwitch13': set(['config'])}
  PROVIDES EventOFPPacketIn TO {'SimpleSwitch13': set(['main'])}
  CONSUMES EventOFPSwitchFeatures
  CONSUMES EventOFPEchoReply
  CONSUMES EventOFPHello
  CONSUMES EventOFPErrormsg
  CONSUMES EventOFPEchoRequest
  CONSUMES EventOFPPortStatus
  CONSUMES EventOFPPortDescStatsReply
connected socket:<eventlet.greenio.base.GreenSocket object at 0xb63c8aec> address:('127.0.0.1', 5500
4)
hello ev <ryu.controller.ofp_event.EventOFPHello object at 0xb63c840c>
move onto config mode
EVENT ofp_event->SimpleSwitch13 EventOFPSwitchFeatures
switch features ev version=0x4,msg_type=0x6,msg_len=0x20,xid=0x92b165c,OFPSwitchFeatures(auxiliary_i
d=0,capabilities=71,datapath_id=1,n_buffers=256,n_tables=254)
move onto main mode
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:01 ff:ff:ff:ff:ff:ff 1
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:02 00:00:00:00:00:01 2
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:01 00:00:00:00:00:02 1
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:01 ff:ff:ff:ff:ff:ff 1
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:03 00:00:00:00:00:01 3
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:01 00:00:00:00:00:03 1
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:02 ff:ff:ff:ff:ff:ff 2
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:03 00:00:00:00:00:02 3
EVENT ofp_event->SimpleSwitch13 EventOFPPacketIn
packet in 1 00:00:00:00:00:02 00:00:00:00:00:03 2
```

Open vSwitch & OpenFlow

Open vSwitch

- A production quality, multilayer **virtual switch**
- It can operate both **as a soft switch** running within the hypervisor, and as the **control stack for switching silicon**.
- It has been ported to **multiple virtualization platforms and switching chipsets**.

Open vSwitch Commands

- Show Interfaces

```
$ sudo ovs-vsctl show
```

- Set Controller on the specific OvS

```
$ sudo ovs-vsctl set-controller <Bridge>:tcp:1.2.3.4.6633
```

- Set OpenFlow version that the specific OvS can support

```
$ sudo ovs-vsctl set bridge <Bridge> protocols=OpenFlow13
```

- Clear the setting of OpenFlow version on the specific OvS

```
$ sudo ovs-vsctl clear bridge <Bridge> protocols
```

- Query the flow rules on the specific OvS

```
$ sudo ovs-ofctl dump-flows <Bridge> -O OpenFlow13
```

Result - Mininet

\$ mininet> pingall

```

*** Stopping 1 switches
s1
*** Stopping 3 hosts
h1 h2 h3
*** Done
completed in 27.992 seconds
mininet@mininet-vm:~$ ryu-manager --verbose ryu.app.simple_switch_13
ryu-manager: command not found
mininet@mininet-vm:~$ sudo mn --topo single,3 --mac --switch ovsk,protocols=OpenFlow13 --controller
remote
*** Creating network
*** Adding controller
Unable to contact the remote controller at 127.0.0.1:6653
Unable to contact the remote controller at 127.0.0.1:6633
Setting remote controller to 127.0.0.1:6653
*** Adding hosts:
h1 h2 h3
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1) (h3, s1)
*** Configuring hosts
h1 h2 h3
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> ryu-manager --verbose ryu.app.simple_switch_13
*** Unknown command: ryu-manager --verbose ryu.app.simple_switch_13
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3
h2 -> h1 h3
h3 -> h1 h2
*** Results: 0% dropped (6/6 received)
mininet>

```


Result - Open vSwitch s1

Flow table (7 rules)

Top 6 rules: host to host

The other: to the controller

```
mininet@mininet-vm:~$ sudo ovs-ofctl dump-flows s1 -O OpenFlow13
OFPST_FLOW reply (OF1.3) (xid=0x2):
  cookie=0x0, duration=23.221s, table=0, n_packets=2, n_bytes=140, priority=1,in_port=2,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:03 actions=output:3
  cookie=0x0, duration=23.228s, table=0, n_packets=3, n_bytes=238, priority=1,in_port=3,dl_src=00:00:00:00:00:03,dl_dst=00:00:00:00:00:01 actions=output:1
  cookie=0x0, duration=23.222s, table=0, n_packets=3, n_bytes=238, priority=1,in_port=3,dl_src=00:00:00:00:00:03,dl_dst=00:00:00:00:00:02 actions=output:2
  cookie=0x0, duration=23.227s, table=0, n_packets=2, n_bytes=140, priority=1,in_port=1,dl_src=00:00:00:00:00:01,dl_dst=00:00:00:00:00:03 actions=output:3
  cookie=0x0, duration=23.232s, table=0, n_packets=2, n_bytes=140, priority=1,in_port=1,dl_src=00:00:00:00:00:01,dl_dst=00:00:00:00:00:02 actions=output:2
  cookie=0x0, duration=23.233s, table=0, n_packets=3, n_bytes=238, priority=1,in_port=2,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:01 actions=output:1
  cookie=0x0, duration=28.92s, table=0, n_packets=9, n_bytes=546, priority=0 actions=CONTROLLER:65535
mininet@mininet-vm:~$
```

A Useful Tool - IPerf

- Create TCP or UDP data stream to measure the throughput of a network
- To become a iperf server

```
$ iperf -s [-u]
```

- To become a iperf client and send data frame

```
$ iperf -c <ip> [-u] [-n <packet_size>] [-b <bandwidth>]
```

- The controller need to detect whether the uploading traffic of a host is beyond limitation.

Reference

- Mininet ([link](#))
- Mininet Walkthrough ([link](#))
- Data center topology ([link](#))
- Open vSwitch Commands ([link](#))
- OpenFlow Messages ([link](#))
- Ryu Book ([link](#))
- Switching Hub ([link](#))
- Spanning tree ([link](#))
- Traffic Monitor ([link](#))
- Topology Viewer ([link](#))